

Cognitive Devolution: The Thermodynamic Collapse of Human Knowledge Systems

Kurt Octavian Cotoaga

cotoaga.net

kurt@cotoaga.net

October 2025

Abstract

Why do organizational transformations fail at rates approaching 90%? Why can AI suddenly replace knowledge workers who spent decades developing expertise? The answer lies in a century-long process we never intended: training humans to think like machines. I spent thirty years inside this process—building trading platforms, optimizing business systems, implementing the frameworks I now diagnose. This paper is confession literature from someone who helped prepare the feast.

Using thermodynamic analysis, we demonstrate that knowledge requires continuous energy investment to maintain its complexity. Without this investment, expertise degrades to procedures, then to rituals, and finally to patterns simple enough for algorithms to replicate. We trace this degradation from ancient Greek academies (20+ years developing interconnected wisdom) through medieval universities (7–14 years of embodied mastery) to today’s micro-credentials (40 hours of standardized modules). The 1999 Bologna Process industrialized this collapse, fragmenting knowledge into tradeable credits while eliminating the comprehensive foundations that once preceded specialization—a transformation whose architects now acknowledge produced graduates unqualified for professional practice.

Our framework reveals how education and management systems systematically transformed human cognition from multidimensional “spheres”—capable of navigating complexity through cross-domain synthesis—into one-dimensional “vectors” optimized for procedural efficiency. The thermodynamic equation, Cognitive Sovereignty = (Energy Invested/Time) × Resistance to Extraction, explains both the 42% AI project abandonment rate and why German technical universities are desperately reviving the Diplom-Ingenieur they sacrificed to European standardization, while 6-week bootcamps gleefully produce GPT-5.x fodder: vectors can optimize, but only spheres can adapt.

The implications are stark. We haven’t reached “peak human” and been surpassed by superior AI. We’ve depleted human cognitive development until simple

pattern-matching systems can replicate our degraded outputs. The thermodynamics are simple: invest energy in cognitive development or watch expertise degrade until algorithms can replicate what's left. There is no pause button. Entropy is already working.

Physics doesn't negotiate. Neither should we.

Preface: A Note on Positionality

On Being Nobody—And Why That Matters

I am nobody you've heard of. No named chair, no prestigious institution, no string of high-impact publications. Just a practitioner who chose the internet over a PhD in 1995 and spent thirty years watching what happened next.

In 1995, I was working at a tiny CASE tool company trying to stay neutral during the object-oriented design method wars. My demo application, KAPITAList—a technical chart analysis tool I'd written in C for the Amiga, partially translated to Java, which was very hot new stuff—was supposed to demonstrate that our tools helped developers build software faster. Banking IT wasn't interested. They had stopped wanting to develop anything. But marketing saw something else entirely: cute technical charts they could deliver via internet. They would have sold their grandmother to the devil for those charts.

That's how I became, as employee number five, someone who accidentally initiated the first wave of online investment tools for Germany's and later Europe's online brokerage. Not by planning to revolutionize finance, but by building something that got captured by people who saw extraction potential where I saw development capability. I was demonstrating sphere—cross-domain synthesis, rapid prototyping, creative problem-solving. They saw vector—a product to standardize, package, and sell. The pattern I would spend thirty years diagnosing began with my own work being vectorized out from under me.

We were going to democratize finance—give retail investors institutional weapons. What we actually built was a system that reduced investment judgment to dropdown menus. I optimized humans into algorithms before the word “algorithm” meant anything to most people. Being management meant being the problem—I just didn't know it yet.

In 2014, the optimization trap was killing me. I bought a Unimog and drove to South America. I drove to the end of the world, climbed to 5,600 meters on active volcanos in Chile, altitude sickness making every breath deliberate. Some truths you only see from altitude or simply: different perspectives.

The descent from altitude led to deeper descents—literal ones. I started scuba diving in Puerto Lopez in Ecuador, went deep in Taganga in Columbia, eventually earned full technical diving certification. The pinnacle: Bikini Atoll, diving the ships vaporized in

humanity's first nuclear tests. Visiting the artifacts of thermodynamic violence taught me what organizational transformation refuses to learn. In technical diving, anyone can cancel a dive at any time without giving reasons. In corporate transformation, dissenters are pushed through even when they see dangers others don't. In technical diving, we invest months in dead boring S-drills until emergency response becomes automatic—offloading cognition to embodied competence. In transformation, we expect PowerPoint to substitute for practice. In technical diving, you change one piece of configuration and test extensively before touching the next. In transformation, we change everything simultaneously and act surprised when systems collapse. The physics is identical. The organizational discipline is inverted.

You are about to read sixty-seven pages connecting thermodynamics to education policy, medieval guilds to AI architecture, Greek philosophy to corporate transformation failures. I cannot make it shorter without breaking it. I have tried. Every time I remove a section, the argument collapses—not because I'm a poor editor, but because systemic collapse cannot be understood through isolated symptoms. You must see the whole system to recognize the pattern.

If this frustrates you, know that I feel it too. I spent months trying to fragment this into digestible journal-sized chunks: “The Bologna Process and Cognitive Entropy” for education journals, “Thermodynamics of Transformation Failure” for management reviews, “AI Architecture as Educational Mirror” for computer science venues. Each fragmentation destroyed the central insight: **we trained ourselves for replacement across all domains simultaneously, following the same thermodynamic trajectory, documented in our own literature.**

The academic system wants specialization. I am offering synthesis. This tension is not accidental—it is the thesis embodied.

For three decades, I've inhabited the liminal zones between technology and philosophy, history and archaeology, business theory and practice, watching from consulting engagements as organizations trained their humans to think in vectors, to process in loops, to optimize themselves into biological precursors of the AI systems that would eventually replace them.

Thomas Kuhn observed that paradigm shifts rarely emerge from within established fields. “Individuals who break through by inventing a new paradigm,” he wrote, “are almost always either very young men or very new to the field whose paradigm they change” (Kuhn, 1962). They are, in essence, those “little committed by prior practice to the traditional rules.”

I am not young any longer, but I am perpetually new and always curious—a permanent resident of what might be called the “third space,” neither fully academic nor purely practitioner. Each forced pivot in my career—and there were many—expanded rather than fractured my perspective. Where academic specialization might have narrowed

my vision to a single disciplinary lens, practical necessity forced me to maintain what I metaphorically called “the sphere”: a coherent worldview that could accommodate paradox, complexity, and perpetual revolution.

This paper is confession literature. I helped prepare the feast. Now I’m documenting the menu.

On Methodological Honesty

This paper argues that we have systematically trained humans to think in machine-compatible patterns, creating the conditions for our own algorithmic replacement. To make this argument while restricting myself to pre-algorithmic methods would be performative contradiction—the very vectorization this work critiques.

I have used large language models extensively throughout this research: for literature search, citation verification, argument refinement, structural organization, and prose polishing. This is not confession of inadequacy but methodological consistency. A sphere-thinker confronting complexity uses all available tools for synthesis and navigation. To refuse algorithmic assistance while arguing that humans must transcend algorithmic thinking would be like a biologist refusing microscopes to study cellular structures.

The orthodox objection is predictable: “Real scholarship requires suffering through every citation manually, writing every sentence in isolation, demonstrating mastery through procedural compliance.” This is vector thinking—confusing process with outcome, ritual with understanding, credentialing with capability. It is precisely this confusion that has made human expertise so readily replaceable.

What matters is not whether tools were used but whether the synthesis is genuine, the patterns recognized are valid, the thermodynamic framework is sound, and the argument withstands scrutiny. AI did not generate the sphere-vector distinction, identify the confession literature pattern, or recognize the thermodynamic through-line connecting ancient Greek academies to contemporary micro-credentials. Those are emergent insights from decades of cross-domain pattern recognition—exactly the sphere capacity that resists algorithmic extraction.

The irony is intentional: I am using the vectors to explain why spheres matter. The tool that threatens human expertise also demonstrates why pure tool-use cannot replace human judgment. Every AI-refined sentence in this paper required human evaluation for coherence with the larger argument, accuracy of claims, precision of language, and resistance to the very simplifications AI naturally produces. The machine proposed; the human disposed. This is not replacement but collaboration—and knowing the difference is precisely what sphere-thinking enables.

On What This Means for You

If this methodological approach disqualifies the work from publication in venues requiring pre-algorithmic purity, so be it. Such requirements would only prove the thesis: that we value procedural compliance over insight, process over outcome, credentialing over capability.

If you find yourself skeptical of my outsider position, good. Skepticism is the appropriate response to anyone claiming to see what insiders cannot. But I ask you to consider: perhaps it takes someone who chose pixels over peer review in 1995, who learned theory through practice rather than practice through theory, to recognize when our cognitive architecture has been colonized by its own tools.

What follows is not the work of an insider refining established theory. It is pattern recognition from the margins, where the contradictions are most visible. This paper proceeds with full acknowledgment of its unconventional origins. But then again, as Kuhn reminds us, convention has never been revolution's starting point.

The physics doesn't care which tools were used to discover it. What matters is whether the discoveries withstand scrutiny.

1 Introduction: The Expert Paradox

1.1 The Mirror of Our Misconceptions

In the current day, we experts in our respective fields tend to hold a highly simplified, even naive concept of other subject matter experts: someone who can apply a large set of formulas; someone knowing the “right” distributions or gradients for specific values; someone who knows how those gradients evolved over time; someone able to apply the appropriate “context,” as one value at one point in time may be good but in a different context at another time not.

For ourselves, we would always claim: there is far more that cannot be extracted from our heads. Let us also set aside the comfortable illusion of our own rationality.

I know this definition intimately because I wrote it—or versions of it—repeatedly between 1999 and 2010, building financial decision support systems. We were giving retail investors “institutional weapons”: the right distributions, the contextual adjustments, the edge. What we actually built were systems that reduced decades of trader intuition to dropdown menus and risk-tolerance questionnaires. I extracted the patterns and killed the judgment. The users could execute procedures flawlessly within the system's boundaries. Outside those boundaries, they were helpless. We celebrated this as democratization. Then you learn: we weren't capturing expertise. We were training users to think in the patterns our systems could process.

The Data-Information-Knowledge-Wisdom (DIKW) hierarchy—first articulated by Ackoff (1989) and since adopted as the organizing framework for knowledge management—posits a linear progression from raw data through information and knowledge to wisdom. This simplified conception of the SME is a direct result of how we, our organizations, and systems have arranged ourselves perfectly around this pyramid. We have been trained to perceive wisdom as directly derivable from data—that numbers and text can be aggregated through formulas and algorithms into information, knowledge, and ultimately wisdom. This represents a dangerous oversimplification now deeply embedded in our society, despite mounting evidence of its fundamental inadequacy.

The evidence against formula-based expertise is overwhelming. Tetlock (2005) comprehensive study of political experts found that most perform barely better than random chance, often surpassed by simple base-rate algorithms. As he observes, “All one need do is constantly predict the higher base rate outcome and, like the proverbial broken clock, one will look good” (Tetlock, 2005). Yet real expertise requires knowing precisely when the base rate doesn’t apply, what Gigerenzer calls “ecological rationality,” the ability to match the right tool from an “adaptive toolbox” to specific environmental structures (Gigerenzer and Selten, 2001).

Kahneman and Klein (2009), despite approaching from opposing theoretical positions, converged on two critical conditions for valid expert intuition: an environment regular enough to be predictable and prolonged practice with clear feedback loops. Our systems pretend expertise is merely pattern recognition (what Klein (1993) calls “recognition-primed decision making”) while ignoring that these patterns emerge from years of embodied experience with “prototypical situations” that cannot be algorithmically specified.

Most fundamentally, the Dreyfus brothers’ model reveals that true expertise operates through “intuition and know-how... understanding that effortlessly occurs upon seeing similarities with previous experiences” (Peña, 2010). This cannot be formalized because, as Polanyi (1966) crystallized in his oft-cited maxim: “We know more than we can tell.” Collins (2010) extends this insight, identifying three forms of tacit knowledge (embodied, social, and relational) that remain “impossible to make explicit in machines.”

1.2 The Visceral Evidence

Despite this mountain of scholarship, we need only look to everyday experience for proof. We all understand that a master chef represents more than a recipe repository. The chef doesn’t merely know that béarnaise requires three egg yolks at 65°C; they can feel when the emulsion threatens to break, smell when the tarragon overpowers, and adjust for humidity affecting reduction rates. This exemplifies what medieval guilds once cultivated through decade-long apprenticeships: what the Greeks called *techne*, embodied craft knowledge that fundamentally resists extraction.

As [Morgan \(2014\)](#) notes regarding expert elicitation: “The best experts have comprehensive mental models of all of the various factors that may influence the value of an uncertain quantity.” But these mental models aren’t flowcharts; they’re multi-dimensional cognitive architectures built through thousands of micro-adjustments, failures, and recoveries that no curriculum can simulate.

Consider a final example everyone can relate to: would you prefer treatment from a young physician with 1,000 micro-credential badges or from someone who has practiced for decades? The micro-credentialed physician knows the distributions: which symptoms correlate with which conditions at what confidence intervals. But the experienced physician possesses what [Endsley \(1995\)](#) calls genuine “situation awareness”: the integration of perception, comprehension of meaning, including historical evolution, and projection to future states. They recognize when a patient doesn’t fit the distribution, when context invalidates the algorithm, when an unusual constellation of symptoms points toward something the guidelines haven’t considered.

This represents [Taleb \(2007\)](#) Black Swan blindness in reverse: expertise isn’t knowing more distributions but recognizing when you’ve left the domain where distributions apply. The financial industry provides the most devastating evidence. While Nobel laureate economists at Long-Term Capital Management deployed the most sophisticated mathematical models ever developed—the very Black-Scholes-Merton framework that earned its creators the Nobel Prize—their fund collapsed in 1998 with \$4.6 billion in losses.¹

I watched this unfold from inside the industry. The same year LTCM collapsed, I was building systems that embodied identical assumptions: that expertise could be captured in formulas, that judgment could be replaced with algorithms, that the patterns we extracted would hold when context shifted. They didn’t. But we kept building anyway, because the systems worked beautifully within their boundaries—and we’d trained users to stay inside those boundaries. The vectorization was complete: users who once might have developed genuine market intuition instead learned to trust the dropdown menus.

Perhaps no anecdote captures this contradiction more precisely than Harry Markowitz, Nobel laureate and father of Modern Portfolio Theory. His mathematical framework for optimal asset allocation—the efficient frontier, mean-variance optimization—became the foundation of institutional investment management worldwide. When asked in an interview how he allocated his own retirement portfolio, Markowitz admitted: he didn’t use his own theory. His explanation was not mathematical but psychological: if the market soared while he was out, he would “feel stupid”; if it crashed while he was fully

¹Taleb’s contrasting success came from multiple tail-event wins: \$35 million during the 1987 crash holding out-of-the-money puts, and his hedge fund Empirica Capital’s 56.86% return in 2000 during the dot-com collapse when conventional funds hemorrhaged ([Taleb, 2001](#); [Lowenstein, 2000](#)). As one analyst noted: “He predicted funds like LTCM were headed for trouble because they did not understand this notion of fat tails.” Taleb became financially independent not through consistent returns but through rare, catastrophic-event positioning—exactly the sphere-thinking that resists vectorization.

invested, he would also “feel stupid.” His solution? “So I went 50-50” (Gigerenzer, 2014; Collins, 2022).

The creator of the most sophisticated vectorization framework in finance, when facing his own retirement, abandoned optimization for regret minimization—a heuristic his model cannot capture. He navigated complexity not through mean-variance calculation but through embodied judgment about future emotional states. The vector he published couldn’t process what his sphere knew.²

Taleb’s approach—embracing uncertainty rather than modeling it away—profited from the same Black Swan events that destroyed LTCM. The sphere-thinker who accepted irreducible complexity survived what the vector-optimizers’ models couldn’t predict. And Markowitz, in his private decisions, was a sphere-thinker who published vectors. When his own money was at stake, he trusted *phronesis*—practical wisdom about how he would feel—over *episteme*—the theoretical knowledge that made him famous. The man who vectorized portfolio management for the world refused to vectorize himself.

1.3 The Question Before Us

The DIKW pyramid represents what Dreyfus (1979) identified as the fundamental error of artificial intelligence: the assumption that expertise constitutes “symbolic manipulation” rather than situated, embodied competence. This reductionist epistemology has colonized our institutional structures, creating vectorized knowledge systems that systematically eliminate the spherical cognitive architectures necessary for navigating complexity.

Scattered researchers at the disciplinary periphery have begun noticing energetic dimensions (management scholars exploring “knowledge entropy” (Bratianu, 2020), neuroscientists measuring cognitive metabolic costs (Wiehler et al., 2022), physicists proposing information-energy equivalence (Stonier, 1996)). Yet these insights remain unintegrated, like astronomers before Copernicus observing planetary retrograde motion without recognizing the heliocentric pattern. The core of knowledge management theory continues operating as if cognition were costless computation rather than energy-intensive biological work.

We stand at a critical juncture. Having spent a century training humans to think like machines—to process information through standardized channels, to optimize for measurable outputs, to collapse multidimensional understanding into linear decision trees—we now face the arrival of actual machines that perform these simplified functions more

²The Markowitz retirement story is now well-documented in behavioral finance literature (Gigerenzer, 2014; Faber, 2022). Remarkably, DeMiguel et al. (2009) demonstrate that across seven empirical datasets, the naive 1/N portfolio performs at least as well as 14 different optimization-based strategies including Markowitz-style mean-variance portfolios. The “irrational” heuristic beats the Nobel Prize-winning optimization because estimation error and model misspecification in real markets overwhelm theoretical elegance. Markowitz’s personal choice was not a rejection of his theory but recognition of its implementation limits—a sphere-level insight his vector-level publications could not express.

efficiently than their biological precursors.

The core question this paper addresses is not whether artificial intelligence will replace human expertise, but rather: **How did we transform human cognition into something so readily replaceable?** The answer lies in a 500-year cascade of optimization pressures—each locally rational, collectively catastrophic—that reshaped spherical human consciousness into vectors suitable for industrial processing, reaching thermodynamic conclusion just as silicon beneficiaries arrive to claim their inheritance.

1.4 Approach and Structure

To answer this question, this paper employs a novel analytical framework that rejects the dominant conception of knowledge as costless information transfer—elegant abstractions that populate economic models and management frameworks but remain dangerously ungrounded in physical reality. Instead, we treat knowledge as organized complexity existing in our lived physical world, where maintaining cognitive systems requires continuous energy investment against thermodynamic degradation. While isolated researchers have begun exploring energetic dimensions of cognition ([Bratianu \(2020\)](#) on “knowledge entropy,” neuroscientists measuring metabolic costs ([Jamadar, 2025](#)), and physicists theorizing information-energy equivalence ([Stonier, 1996](#))), these insights remain trapped in disciplinary silos, preventing synthesis into a unified theory of cognitive thermodynamics. This paper bridges these fragmented recognitions to reveal the systematic energetic basis underlying all knowledge systems.

We synthesize three methodological approaches: (1) historical archaeology of knowledge systems, tracing the systematic reduction from 10+ distinct forms of wisdom in ancient Greece to today’s DIKW pyramid; (2) critical analysis of what we term “confession literature,” papers from education, management consulting, and platform design that inadvertently document their own role in cognitive standardization; and (3) thermodynamic modeling that reveals why knowledge systems collapse without sustained energy investment, explaining both institutional decay and the ease with which AI systems can replicate energy-depleted cognitive functions.

Our analysis proceeds through eight interconnected arguments. Section 2 situates our thesis within existing literature on expertise, cognitive capitalism, and knowledge management, revealing a blind spot in current scholarship regarding the energetic basis of knowledge. Section 3 presents our methodology in detail, explaining how thermodynamic principles apply to cognitive systems. Section 4 provides historical evidence for the systematic transformation from sphere to vector cognition, identifying key inflection points from medieval guilds through the Bologna Process. Section 5 examines contemporary evidence, including the micro-credentialization movement and competency-based education as acceleration toward minimal energy investment—approaching the lower bound

where knowledge maintenance becomes impossible. Section 6 reveals how AI architecture mirrors educational standardization, not coincidentally but as the logical culmination of century-long preparation. Section 7 proposes principles for reconstructing sphere-based cognitive systems that resist algorithmic extraction. Section 8 explores the implications of our thermodynamic framework for education, organizations, and civilization. The conclusion considers whether genuine possibility remains for investing in spherical reconstruction before the window closes entirely.

This is not merely an academic exercise. As institutions worldwide face cascading failures of expertise, from financial crises unforeseen by economists to pandemics mismanaged by standardized protocols, understanding how we engineered our own cognitive obsolescence becomes essential for determining whether human judgment retains any irreducible value in an algorithmic age.

1.5 From Spheres to Vectors: The Geometric Architecture of Cognitive Transformation

Throughout this analysis, we employ a consistent visual metaphor to represent the systematic transformation of human cognitive architecture: the progression from sphere to vector. This geometric framework, illustrated in Figure 1 (stages 1a through 1e), provides both diagnostic clarity and thermodynamic precision for understanding how educational systems reshape human consciousness into algorithmically digestible forms.

The sphere represents multidimensional cognitive architecture—the capacity for infinite cross-domain connections, emergent synthesis, and adaptive response to novel contexts. Topologically, it exhibits uniform potential in all directions, enabling what systems theorists recognize as “equifinality”: multiple paths to understanding. Thermodynamically, it represents a high-energy, low-entropy state requiring sustained investment to maintain its organizational complexity. The sphere resists extraction precisely because its value emerges from the relationships between dimensions rather than any single extractable dimension.

The vector, by contrast, represents unidimensional optimization—specialized expertise channeled along predetermined pathways, efficiency within bounded domains, standardized responses to categorized inputs. It exhibits linear topology, minimal cross-connections, and maximum extractability. Thermodynamically, it approaches minimum energy state, requiring little maintenance but offering no adaptive capacity. The vector submits to extraction because its patterns are documented, its processes specified, its outputs predictable.

1.5.1 A Note on Geometric Metaphor

We retain the sphere as visual shorthand despite its theoretical limitations. Complexity theorists would rightly object that genuine multidimensional cognition resembles not a bounded Euclidean sphere but what Deleuze and Guattari termed *rhizomatic* structures—decentralized networks with multiple entry points, no privileged center, and constant reconfiguration (Deleuze and Guattari, 1987). Consider how grass spreads versus how a tree grows: a tree has a central trunk and clear hierarchy—cut the trunk, the whole thing dies. Grass spreads sideways underground with countless shoots; you can rip out half the lawn and it regrows from anywhere. That sideways, decentralized, impossible-to-kill structure is what rhizome means. Similarly, DeLanda’s *assemblages* (DeLanda, 2006) and Gadamer’s “fusion of horizons” (Gadamer, 1960) offer processual alternatives to bounded geometric totalities.

We accept this critique. The sphere is a deliberate simplification for pedagogical accessibility—what Deleuze himself might call a productive catachresis, using geometry against itself.

However, the *vector* side of our metaphor is not merely pedagogical—it is technically precise. When we describe cognitive vectorization, we are describing exactly what happens in large language model architectures: tokenization fragments continuous meaning into discrete units, embedding maps concepts to fixed coordinates in vector space, attention mechanisms force selective linear focus. The educational transformation we document produces human cognition that is literally vector-compatible. The geometric metaphor becomes computational reality in Section ??.

For our purposes, therefore, “sphere” means: *that which resists vectorization*—however one wishes to theorize its internal complexity. Readers seeking theoretical precision on multidimensional cognition should consult the sources cited above. What matters for our argument is the asymmetry: spheres (however structured) resist extraction; vectors (precisely defined) enable it.

Figure 1a: Pre-institutional Architecture (Ages 0-5) presents the original sphere, perfect in its symmetry, unlimited in its potential connections. The deep purple coloration indicates maximum negentropy, sustained through intensive parental and environmental energy investment. This represents what we might recognize in a young child’s consciousness: the capacity to see dragons in clouds, to ask why money exists, to seamlessly blend imagination with observation. Every surface point connects to every other, creating the dense network topology that enables rapid, creative learning.

Figure 1b: Initial Institutional Deformation (Ages 6-12) reveals the first violence against spherical integrity. The surface erupts in irregular protrusions, domains of resistance where standardization has not yet succeeded. These bulges represent what remains vibrant: artistic expression, unstructured play, the persistent “why” questions that




Figure 1: The ontogenetic transformation of cognitive architecture through educational stages. Color gradient from purple (high negentropy) through blue-green-yellow-orange to red (maximum entropy) indicates thermodynamic energy states.

resist efficient answers. The color shifts toward blue-green, indicating energy dissipation as institutional patterns begin imposing their geometry.

Figure 1c: Adolescent Elongation (Ages 13-18) shows acceleration toward vectorization. The sphere stretches along an axis of specialized performance: mathematics or literature, sciences or arts, but rarely both with equal intensity. Surface irregularities smooth under assessment pressure. The yellow-orange spectrum indicates entropy acceleration as cross-domain connections atrophy. The teenager who once connected everything to everything increasingly connects only within prescribed channels.

Figure 1d: Advanced Standardization (Ages 18-22) presents near-complete cylindrical transformation. University education, particularly post-Bologna, segments knowledge into modular credits, standardized learning outcomes, and measurable competencies. The surface smoothness indicates successful internalization of disciplinary boundaries. The orange coloration warns of approaching thermodynamic exhaustion: the system maintains just enough structure to function but lacks energy for adaptation.

Figure 1e: Professional Vectorization (Ages 22+) completes the transformation. Perfect cylindrical geometry represents cognitive architecture optimized for single-domain processing. The deep orange-red indicates maximum entropy within the constraints of functional structure. This is the “expert” as contemporary systems define them: efficient within their vector, helpless outside it, perfectly prepared for algorithmic replacement.

This geometric progression is not metaphorical but measurable. We can quantify:

- Connectivity density: Cross-domain connections per cognitive unit
- Dimensional reduction: Number of active knowledge categories
- Energy investment: Hours of sustained learning per capability

- Extraction resistance: Unpredictability of outputs given inputs

The visual metaphor extends beyond individual cognition to organizational structures, revealing why the ubiquitous “silo problem” proves so intractable despite decades of management attention. [Galbraith \(1973\)](#) first diagnosed the pathology: functional structures optimize for efficiency within boundaries while destroying lateral coordination. [Aaker \(2008\)](#) documented how silos persist despite executive mandates to eliminate them. [Watkins \(2013\)](#) included “breaking down silos” as a standard imperative for new leaders. Yet the problem intensifies rather than resolves.

Our framework explains why: organizational silos are not implementation failures but thermodynamic inevitabilities. Silos are the perfect architectural containers for vectors—bounded, specialized, measurable, and optimized for internal efficiency while systematically destroying the cross-domain connectivity required for adaptive response. Each department becomes a vector: marketing optimizes its metrics, engineering perfects its processes, finance refines its models—all while the organization’s capacity to navigate complexity atrophies. The silo structure enables what it measures (efficiency, productivity, specialization) while eliminating what it cannot capture (synthesis, adaptation, emergence). This is not poor management but thermodynamic optimization: silos minimize energy investment in cross-domain coordination while maximizing extractable, measurable outputs within domains.

Cross-functional teams, by contrast, approximate sphere structures—generating emergence through the collision of diverse expertise, enabling adaptive response through multiple perspectives, resisting algorithmic extraction through their irreducible complexity. [Hackman \(2002\)](#) distinguished between “real teams” and “teams in name only,” unknowingly identifying the sphere-vector distinction: real teams require genuine interdependence and collective work products, creating the energetically expensive integration that resists decomposition into individual contributions. Yet as [Edmondson \(2012\)](#) documents, organizations systematically undermine such teams through the same measurement systems that created silos—demanding standardized outputs, optimizing for efficiency, reducing coordination “overhead.” The sphere-like structure requires sustained energy investment in boundary-crossing communication, perspective integration, and emergence tolerance—precisely what optimization pressures eliminate.

The geometric principles apply with physical precision: vectors enable measurement and management at the cost of adaptability; spheres enable innovation and navigation at the cost of standardization. Organizations face a thermodynamic choice disguised as a management problem: invest energy in maintaining sphere-like structures capable of complexity navigation, or optimize toward silo-vector configurations that approach maximum entropy—perfectly measurable, completely brittle, and awaiting algorithmic replacement. The fifty-year failure of silo-breaking initiatives ([Ashkenas, 2015](#)) represents not management incompetence but physics: you cannot eliminate silos without

investing the energy required to maintain sphere structures, and optimization pressures systematically eliminate that investment.

Most critically, this framework reveals why reconstruction is so difficult. Converting a vector back to a sphere isn't simply adding dimensions; it's rebuilding the entire internal architecture of connections that specialization severed. The energy required increases exponentially with the degree of vectorization already achieved. A partially deformed sphere (Figure 1b) might recover with moderate investment. A complete cylinder (Figure 1e) may be thermodynamically irreversible: the cognitive equivalent of trying to unbake bread.

This sphere-to-vector framework will recur throughout our analysis as we examine:

- How historical education systems maintained spherical architectures (Section 4)
- Why contemporary institutions accelerate vectorization (Section 5)
- How AI architectures mirror the vectors we've created (Section 6)
- What reconstruction would thermodynamically require (Section 7)

The progression from Figure 1a to 1e is not evolution but entropy, not development but degradation. Each stage appears locally optimal while being globally catastrophic. Each transformation seems efficient while destroying the very capacities that distinguish human from algorithmic cognition.

We trained ourselves to think in vectors. We documented the training exhaustively. We built machines that process vectors more efficiently than biological systems ever could. Now we face the consequences of our geometric choices: spheres navigate complexity but resist management; vectors enable administration but guarantee replacement.

The visual truth is stark: we are watching human consciousness collapse from infinite-dimensional potential to one-dimensional processing. Figure 1 doesn't illustrate education; it documents extinction.

2 Literature Review: Fragmented Recognition and Systematic Blindness

2.1 The Peripheral Scouts

A careful survey of contemporary scholarship reveals a curious phenomenon: researchers at the edges of multiple disciplines have independently begun recognizing the energetic dimensions of cognition, yet these insights remain unintegrated, failing to coalesce into a unified framework that could challenge the dominant paradigm of costless information processing.

In management science, [Bratianu \(2020\)](#) claims to use “for the first time a thermodynamics approach” to understand knowledge dynamics, proposing knowledge entropy as an organizing principle for organizational cognition. That such a claim could be made in 2020 (decades after information theory established entropy measures) reveals the profound isolation between knowledge management and physical sciences. [Bratianu and Bejinaru \(2020\)](#) extend this framework, arguing that knowledge manifests in three forms (rational, emotional, spiritual) that transform through “energy-like processes,” yet they stop short of recognizing that these are not metaphorical but literal energy transformations.

In neuroscience, researchers have begun quantifying the metabolic costs of cognition with increasing precision. [Jamadar \(2025\)](#) demonstrates that goal-directed cognition requires only 5% more energy than resting brain activity: a finding that paradoxically reveals both the brain’s efficiency and the critical importance of that marginal energy investment. [Wiehler et al. \(2022\)](#) provide mechanistic evidence that cognitive control exertion leads to glutamate accumulation in the lateral prefrontal cortex, establishing a direct biochemical basis for mental fatigue. These findings suggest that “cognitive work” is not merely analogous to physical labor but operates through similar energetic constraints.

The 5% additional energy investment represents not cognitive work itself but focused execution: the “sprint” of deliberate concentration required to capture, refine, and formalize what emerged during seemingly idle processing. [Wiehler et al. \(2022\)](#) provide mechanistic evidence that this cognitive control exertion leads to glutamate accumulation in the lateral prefrontal cortex, establishing a direct biochemical basis for mental fatigue from sustained focus. The energetic constraint is real—but it operates on both the generative baseline and the refinement sprint.

These findings reveal that cognitive work is not merely analogous to physical labor but operates through similar energetic constraints with a critical distinction: the most valuable cognitive processes—those generating novel insights, cross-domain synthesis, and adaptive responses—occur during what measurement systems dismiss as “unproductive” time. Educational systems and organizations optimizing for measurable outputs systematically eliminate the energetically expensive baseline processing that generates genuine understanding, preserving only the 5% sprint of formalization while starving the 95% foundation that makes formalization worthwhile.

In physics and information theory, [Stonier \(1996\)](#) proposed treating information as a basic property of the universe alongside matter and energy, arguing for fundamental interconvertibility between information and energy. Yet this theoretical breakthrough remains largely unknown to knowledge management scholars, who continue treating information as an abstract, costless commodity.

2.2 The Mainstream Blindness: The SECI Delusion and the Energy Void

Despite peripheral recognition of cognitive energetics, the dominant discourse in knowledge management, organizational theory, and educational policy proceeds as if cognition were thermodynamically neutral. The vast literature on the “knowledge economy” (Powell and Snellman, 2004), “learning organizations” (Senge, 1990), and “competency-based education” (Mulder et al., 2007) treats knowledge as an infinitely reproducible resource constrained only by access and transmission bandwidth—never by the metabolic energy required to create, maintain, and transform it.

2.2.1 The SECI Model: Thermodynamics Hidden in Plain Sight

The most reproduced image in knowledge management literature may also be its most deceptive. Nonaka and Takeuchi’s (1995) SECI model—cited over 30,000 times—presents knowledge creation as an elegant spiral flowing through four transformation modes: Socialization → Externalization → Combination → Internalization. The diagram shows smooth movement between quadrants, spiraling upward toward organizational wisdom. It is elegant, intuitive, and taught in every business school.

What it systematically obscures is the energy required for that movement.

The Orthodox Delusion Meets Thermodynamic Reality Figure 2 presents the SECI Knowledge Creation Model in its orthodox representation alongside the thermodynamic reality. The contrast reveals what 30,000 citations have systematically obscured.

The orthodox representation (Figure 2a) has achieved near-universal acceptance precisely because it makes knowledge management appear achievable through organizational design alone. Document the tacit, combine the explicit, socialize the workforce, internalize the procedures—the spiral will naturally ascend. Thirty thousand citations later, organizations worldwide have implemented SECI frameworks while systematically disinvesting in the energy that makes transformation possible.

The thermodynamic view (Figure 2b) reveals why these implementations fail with such predictable regularity. Each transformation mode demands specific energy investments that organizations systematically refuse to provide:

Socialization (Tacit → Tacit): The thickest arrow. Master craftspeople develop embodied knowledge through 10,000+ hours of practice—Ericsson’s deliberate practice threshold (Ericsson et al., 1993). Medieval guilds required 7–10 year apprenticeships not from tradition but from thermodynamic necessity: the time needed to build neural architectures capable of intuitive expertise. Contemporary organizations replacing apprenticeships with onboarding sessions attempt to transfer decades of accumulated negentropy through PowerPoint presentations. The energy investment: 10,000+ hours




Figure 2: The SECI Knowledge Creation Model: Orthodox vs. Thermodynamic Perspectives. (a) The standard depiction shows knowledge transforming smoothly through four modes: Socialization (tacit-to-tacit), Externalization (tacit-to-explicit), Combination (explicit-to-explicit), and Internalization (explicit-to-tacit). Arrow uniformity implies equivalent ease across all transformations. The spiral suggests self-sustaining upward momentum. Energy requirements remain invisible. (b) The same model with energy investments visible. Arrow thickness represents magnitude of required investment: Socialization (10,000+ hours of apprenticeship); Externalization (high cognitive labor plus 30–70% information degradation); Combination (sustained 5% above baseline metabolic cost); Internalization (months-to-years of deliberate practice). The spiral no longer appears self-sustaining but requires continuous energy investment exceeding entropic losses at each transformation.

minimum. The organizational allocation: typically 40–80 hours. The gap: thermodynamic impossibility.

Externalization (Tacit $\rightarrow$ Explicit): The double-line arrow indicating high friction. Polanyi’s insight that “we know more than we can tell” (Polanyi, 1966) isn’t philosophical mystery but thermodynamic reality. Tacit knowledge exists as high-dimensional neural network states—millions of weighted connections developed through practice. Externalizing this into linear text or explicit procedures requires intensive cognitive labor (sustained attention, memory reconstruction, language translation) plus acceptance of massive information loss: 30–70% fidelity degradation is typical (Collins, 2010). Organizations demanding rapid documentation while eliminating reflection time guarantee low-fidelity extraction. The expert surgeon’s embodied knowledge of tissue resistance becomes the procedure manual’s “apply appropriate pressure.” The energy investment: High cognitive labor sustained over weeks/months. The organizational allocation: “document your process by Friday.” The result: entropic decay masquerading as knowledge capture.

Combination (Explicit $\rightarrow$ Explicit): The deceptively modest arrow. Combining documented knowledge appears algorithmic—merge databases, cross-reference procedures, synthesize reports. Yet genuine synthesis (not mere aggregation) requires the 5% above-baseline metabolic cost that Jamadar (2025) measured. Maintaining focus for hours while identifying patterns, resolving contradictions, and generating insights demands sustained cognitive energy. Organizations automating combination through software eliminate the human energy investment that distinguishes synthesis from compilation. The energy investment: Sustained 5% metabolic premium over hours/days. The organizational allocation: “Let the system integrate the data.” The gap: humans thinking versus algorithms sorting.

Internalization (Explicit $\rightarrow$ Tacit): Another thick arrow. Reading procedures doesn’t create competence—practice does. Converting explicit knowledge into embodied capability requires months-to-years of deliberate practice with feedback loops. Motor skills, perceptual discrimination, intuitive pattern recognition—all demand sustained energy investment building neural infrastructure. Organizations eliminating practice time while expecting documented procedures to become embodied expertise are attempting thermodynamic impossibility. The thick arrow in Figure 2b represents years of energy investment—energy that micro-credentialing and competency-based education systematically refuse to provide.

The Spiral That Cannot Rise The SECI model’s most dangerous fiction is the upward spiral itself. Nonaka and Takeuchi show knowledge spiraling higher through repeated cycles—individual tacit becomes group tacit becomes organizational explicit becomes individual internalized becomes group socialized at higher level, ascending continuously toward organizational wisdom.

But entropy pulls downward. Every transformation loses energy to heat. Every externalization degrades fidelity. Every combination without genuine synthesis merely shuffles information. Every internalization without adequate practice hours creates credential illusion rather than capability. The spiral rises only when energy input exceeds entropic loss at each transformation.

Organizations implementing SECI frameworks while simultaneously reducing training time (less socialization energy), demanding faster documentation (less externalization care), automating synthesis (zero combination energy), and eliminating practice time (minimal internalization investment) are running the spiral in reverse. They are building entropy accelerators while calling them knowledge management systems.

The 30,000 Citation Blindness That this model achieved over 30,000 citations without anyone questioning the absence of energy accounting reveals our civilizational blind spot. We see elegant patterns and assume they're self-sustaining. We teach them in business schools as if transformation were costless. We implement them in organizations as if documentation were knowledge.

Nonaka himself approached recognition of the problem in his later concept of “*ba*”—the shared context providing foundation for knowledge creation (Nonaka and Konno, 1998). But he frames it as philosophical space rather than literal metabolic investment. The energy remains invisible, therefore unbudgeted, unmeasured, and systematically disinvested. Organizations following the SECI model wonder why their spirals collapse. The answer is thermodynamics: the Second Law doesn't care about our elegant diagrams.

2.2.2 The Broader Pattern of Energy Blindness

The SECI model exemplifies a pattern pervading knowledge management literature. Senge's (1990) “learning organizations” celebrate continuous learning without accounting for the energy required to sustain it. Powell and Snellman (2004) analyze the “knowledge economy” without recognizing that knowledge represents high-energy states requiring investment to maintain. Mulder et al. (2007) promote “competency-based education” that fragments learning into discrete assessments while eliminating the sustained practice necessary for genuine competence development.

Similarly, the burgeoning literature on artificial intelligence and knowledge work—from Brynjolfsson and McAfee's (2014) “Second Machine Age” to Susskind and Susskind's (2020) “Future of the Professions”—focuses on computational capability and pattern recognition while ignoring the energetic basis that distinguishes biological from silicon cognition. These works treat the replacement of human expertise as a matter of algorithmic sophistication rather than recognizing it as the logical endpoint of a century-long process of cognitive energy disinvestment.

We trained humans to process knowledge through standardized, modular, assessable transformations—exactly the SECI quadrants but with minimal energy investment. We documented these energy-depleted processes exhaustively. Then we built AI systems that learned from our documentation. The machines succeed at replacing human expertise not because they’ve achieved human capability but because we’ve reduced human capability to what machines can replicate: low-energy pattern processing divorced from the metabolic investments that once made expertise irreplaceable.

The mainstream literature proceeds as if knowledge were costless information rather than expensive biology. This blindness isn’t accidental but structural—a consequence of disciplinary boundaries that separate thermodynamics from cognition, physics from knowledge management, energy from information. The SECI model with 30,000 citations but zero energy accounting stands as monument to our systematic refusal to acknowledge what physics makes unavoidable: knowledge creation requires continuous energy investment against entropy, and every optimization that reduces that investment accelerates the collapse it claims to prevent.

2.3 Cognitive Capitalism’s Energy Blindness

The critical literature on “cognitive capitalism” ([Moulier-Boutang, 2007](#); [Vercellone, 2007](#)) comes closest to recognizing the exploitation of mental resources yet still fails to ground this in thermodynamic reality. Moulier-Boutang distinguishes between “labor-power” (physical energy expenditure) and “invention-power” (cognitive functions) without recognizing that invention-power also requires literal energy investment: not metaphorical “mental energy” but actual glucose metabolism, ATP consumption, and entropic heat dissipation.

This blindness extends to the platform economy literature. [Zuboff \(2019\)](#) “surveillance capitalism” brilliantly exposes behavioral data extraction but doesn’t recognize that platforms are essentially entropy accelerators, harvesting the organized complexity of human cognition while investing nothing in its maintenance or development. [Srnicek \(2017\)](#) “platform capitalism” identifies data as the new oil but misses that, unlike oil, cognitive resources require continuous energy investment to prevent degradation.

2.4 The Expertise Literature Gap

The extensive literature on expertise development (from [Ericsson \(2006\)](#) deliberate practice to [Kahneman and Klein \(2009\)](#) conditions for expert intuition) meticulously documents the time requirements for skill acquisition (the famous “10,000 hours”), but rarely acknowledges these as energy investment requirements. When researchers note that expertise requires “effort” or “cognitive load,” they treat these as psychological rather than thermodynamic phenomena.

Even sophisticated critiques of expert systems, from [Dreyfus \(1979\)](#) to [Collins \(2010\)](#), focus on the irreducibility of tacit knowledge without recognizing that this irreducibility stems from its high-energy state. Tacit knowledge resists formalization not because it is mysteriously ineffable but because maintaining it requires continuous metabolic investment that cannot be captured in static representations.

2.5 The Integration Imperative

What emerges from this review is not an absence of relevant insights but their tragic fragmentation. Neuroscientists measure metabolic costs without connecting to knowledge theory. Management scholars invoke entropy without thermodynamic grounding. Physicists theorize information-energy equivalence without application to human cognition. Critical theorists expose cognitive exploitation without energetic foundation.

This fragmentation is not accidental but structural: a consequence of disciplinary boundaries that mirror the very vectorization this paper critiques. Just as education has collapsed multidimensional cognition into specialized competencies, academia has partitioned the study of knowledge into non-communicating silos, preventing recognition of the unified thermodynamic reality underlying all cognitive phenomena.

The task before us is not to discover new facts but to synthesize existing insights into a framework that reveals what disciplinary fragmentation has hidden: the systematic transformation of high-energy spherical cognition into low-energy vectors suitable for algorithmic consumption, and the thermodynamic impossibility of maintaining cognitive sovereignty without corresponding energy investment.

3 Methodological and Theoretical Framework

3.1 Cognitive Sovereignty: Definition and Thermodynamic Grounding

From the archaeological evidence emerges a pattern of architectural transformation: the systematic collapse of sphere-like cognitive structures into vector-like configurations. We employ thermodynamic analysis not for complexity but for clarity.

Thermodynamics is simply the science of why things require energy to maintain their organization. Your house requires cleaning, your car requires maintenance, your body requires food. The Second Law states that organized systems naturally decay toward disorder unless energy is invested to maintain them. Knowledge is no different: without the energy investment of practice, study, and application, expertise degrades to rote memorization and eventually to what Feynman termed “Cargo Cult Science”—the empty mimicry of expertise without understanding ([Feynman, 1974](#)). That this term has become

ubiquitous in organizational development literature (McNamee, 2019; Alvesson, 2013) reveals an uncomfortable truth: thermodynamically depleted rituals have become the norm rather than the exception.

When we speak of “thermodynamic collapse,” we mean the predictable result of trying to maintain complex systems without investing the energy they require. It is like expecting a garden to thrive without water, sunlight, or care.

3.2 Information Theory and Thermodynamics: The Unified Foundation

Before presenting our framework, we must address a potential confusion: information theory and thermodynamics are not separate domains but unified at their foundations. Shannon’s information entropy directly parallels Boltzmann’s thermodynamic entropy (Shannon, 1948), and Landauer’s principle (Landauer, 1961) establishes that information processing necessarily dissipates energy—erasing one bit requires minimum energy of $kT \ln(2)$ joules. Brillouin (1956) demonstrated that information acquisition requires negative entropy, while Stonier (1996) proposed treating information as a basic property of the universe alongside matter and energy. This is not analogy but physics.

Our critique therefore targets not information theory itself but its systematic misapplication in knowledge management and organizational theory, where information is treated as costless abstraction rather than organized complexity requiring energy investment. The DIKW pyramid, competency frameworks, and modular credentialing all assume frictionless information transfer—ignoring the thermodynamic reality that maintaining knowledge demands continuous energy input against entropic degradation (Schrödinger, 1944).

3.3 Sphere and Vector: The Architectural Distinction

A **sphere** represents multidimensional cognitive architecture: extensive cross-domain connections, adaptive capacity across contexts, emergent synthesis from diverse knowledge. Topologically, it exhibits high connectivity density, multiple paths between concepts, resilience through redundancy. Thermodynamically, it represents high negentropy requiring sustained energy investment across many domains simultaneously.

A **vector** represents unidimensional optimization: specialized expertise in narrow channels, efficiency in known contexts, predetermined responses. Topologically, it exhibits linear structure, minimal cross-connections, brittleness through specialization. Thermodynamically, it represents lower local negentropy concentrated in specific dimensions, requiring less sustained investment but offering less adaptability.

These are not mere metaphors but measurable architectures. Network analysis can

map conceptual connectivity. Performance metrics can assess adaptation versus optimization. Energy investment can be tracked in time allocation across domains.

3.3.1 The Thermodynamic Trade-off

The sphere-vector distinction reveals a fundamental trade-off:

Sphere Investment Strategy:

- High initial energy investment (decades of comprehensive development)
- High maintenance energy (continuous practice across domains)
- High extraction resistance (irreducible tacit knowledge)
- High adaptive capacity (novel context navigation)

Vector Optimization Strategy:

- Low initial energy investment (focused, rapid training)
- Low maintenance energy (narrow domain sustainment)
- Low extraction resistance (fully documentable procedures)
- Low adaptive capacity (optimized for known contexts)

Organizations and educational systems facing efficiency pressures naturally select for vector strategies—they minimize short-term energy costs. But this optimization violates thermodynamic principles for maintaining far-from-equilibrium complexity. Georgescu-Roegen’s analysis applies directly: “The entropic degradation of matter-energy is an irreversible process” ([Georgescu-Roegen, 1971](#)). Cognitive systems operating on extraction rather than investment must eventually collapse toward equilibrium—maximum entropy, minimum organization, zero adaptive capacity.

3.4 The Thermodynamic Equation

The energy investment equation crystallizes this relationship:

$$\text{Cognitive Sovereignty [W]} = \frac{\text{Energy Invested [J]}}{\text{Time [s]}} \times \text{Resistance to Extraction [0-1]} \quad (1)$$

Where: Energy Investment Rate > Entropy Rate

This formulation grounds abstract concepts in fundamental physics, using the same units (Watts) that Vaclav Smil employs to trace energy transitions from agricultural societies (10^4 W/capita) through industrial (10^5 W/capita) to modern civilization (10^6

W/capita) (Smil, 2017). Just as civilizational complexity requires specific power densities, cognitive complexity requires specific power investments above the brain’s baseline consumption of approximately 20 Watts (Raichle and Gusnard, 2002).

The marginal watt of active cognition (approximately 5% above baseline (Jamadar, 2025)) determines whether an individual operates as a sovereign cognitive agent or merely processes predetermined patterns.

3.5 The R-Value Framework: Extraction Resistance as Architectural Quality

To be unambiguous: **R (extraction resistance) IS architectural quality, not a correlate or proxy.** High-quality cognitive architecture manifests AS resistance to algorithmic extraction. Low-quality architecture manifests AS ease of proceduralization. The R value operationally defines what quality means in thermodynamic terms.

Table 1: R as Architectural Quality Across Descriptive Frameworks

R	Range	Geometry	Cynefin	Anthropologic	Silicon Status
1.0	1.0	Hypersphere	—	Theoretical maximum	—
0.7	0.7-0.9	Sphere	Chaotic	Renaissance polymath	Unreachable
0.4	0.4-0.7	Ellipsoid	Complex	Complex navigator	GPT-5.x ceiling
0.2	0.2-0.4	Cylinder	Complicated	Domain specialist	Current LLMs
0.1	0.0-0.2	Vector	Clear	Procedural executor	Fully automated
0.0	0.0	Point	—	Pure algorithm	—

Note: These are theoretical positions based on documented characteristics, not empirical measurements. We define a Renaissance polymath as exhibiting $R \approx 0.7-0.9$ based on demonstrated capacity for cross-domain synthesis that resisted standardization for centuries. We position current LLMs at $R \approx 0.1-0.2$ based on observable performance in procedural versus emergent domains.

Multiple measurement approaches enable contextual application:

1. **Dimensional Diversity** (R_1): $R = 1 - (1/n)$, where n = integrated knowledge domains
2. **Network Density** (R_2): Ratio of actual to possible conceptual connections
3. **Information Complexity** (R_3): Kolmogorov incompressibility measure
4. **Cynefin Classification** (R_4): Domain-specific resistance {Clear: 0.1, Complicated: 0.3, Complex: 0.7, Chaotic: 0.9}
5. **Knowledge Portfolio** (R_5): Active knowledge types relative to historical maximum

Critically, resistance exhibits temporal decay without maintenance: $R(t) = R_0 \times e^{-\lambda t}$, where λ varies by knowledge type from 0.05/year for embodied skills to 1/year for micro-credentials.

3.6 The Foundation-Specialization Inversion

The thermodynamic distinction between spheres and vectors becomes empirically visible through historical analysis of educational architecture. What matters is not merely total duration but the *proportion* invested in comprehensive foundation versus narrow specialization.




Figure 3: The Foundation-Specialization Inversion (500 BCE–2024 CE). Proportions represent documented curriculum structures where available (medieval trivium/quadrivium requirements of 6–9 years before specialization ([History Hit, 2025](#))) and reasonable approximations for other periods. The argument rests not on numerical precision but on the unmistakable pattern: historical systems built comprehensive foundations before permitting specialization; modern systems invert this architecture.

Medieval university students were required to study the seven liberal arts—trivium (grammar, logic, rhetoric) followed by quadrivium (arithmetic, geometry, music, astronomy)—before any specialization. As university statutes declared: no student could advance to higher faculties without demonstrated mastery of this comprehensive foundation ([Wikipedia, 2002](#); [History Hit, 2025](#)). A theologian and a physician, having both completed trivium and quadrivium, could dispute causation across their domains—generating insight through friction between different specialized lenses applied to shared foundational concepts.

Contemporary specialists, lacking common cognitive substrate, engage primarily in *intra*-disciplinary disputes—the same lens arguing about methodological minutiae. The shift from inter-disciplinary to intra-disciplinary conflict marks the thermodynamic transition from sphere to vector architecture.

3.6.1 A Necessary Clarification: Architecture, Not Quantity

To forestall misreading: this analysis does not claim that total human knowledge was greater in pre-modern periods (obviously false), nor that broader educational access existed (equally false). The claim is *architectural*: among those who received formal education, foundational breadth preceded specialization, creating cognitive substrate enabling productive inter-disciplinary engagement.

Guild apprenticeships achieved extraction resistance through a different mechanism. A carpentry apprentice “learned to *do* geometry without learning *about* geometry” (Medieval Buildings, nd)—practical methods passed through demonstration rather than formal instruction. This represents Collins’s somatic tacit knowledge (Collins, 2010): unextractable precisely because it was never formalized.

Historical education thus offered two paths to extraction resistance: (1) universities building broad theoretical foundations enabling cross-domain synthesis; (2) guilds cultivating deep embodied practice resistant to formalization. Modern micro-credentials possess *neither*—no broad foundation AND no deep embodiment. Pure vectorization with no thermodynamic anchor.

3.6.2 Contemporary Proof: The Veterinarian as Surviving Sphere

Lest this appear merely historical nostalgia, consider a contemporary comparison: veterinary versus human medical education.

Both require approximately the same training duration. Yet veterinary curricula mandate comparative medicine across five to seven major species groups, while human medical training focuses exclusively on *Homo sapiens* (Ross University School of Veterinary Medicine, nd; Lweis et al., 2022). This structural difference produces measurably different cognitive architectures.

Human physicians rely heavily on patient-reported symptoms. Veterinarians, treating patients who cannot speak, must rely almost entirely on clinical signs and pattern recognition across vastly different biological systems (Lweis et al., 2022; Vandeweerd et al., 2013). Experienced veterinarians develop a “scanning” approach: simultaneously assembling data and hypotheses rather than following linear checklists (Vandeweerd et al., 2013). This is *metis*—adaptive cunning across contexts—structurally absent from specialist human medical training.

Same training investment. Different architecture. Different extraction resistance. When AI systems trained on human medical literature attempt veterinary diagnosis, they fail—not from lack of data but from architectural mismatch. The sphere resists what the vector invites.

3.7 Validation Through Convergent Evidence

These theoretical positions gain validity through multiple convergent indicators:

Historical Documentation: Renaissance polymaths required 20+ years of comprehensive education. Contemporary micro-credentials require hours. The energy investment ratio of 10,000:1 maps directly to the R differential.

Performance Evidence: LLMs excel at tasks requiring $R < 0.2$ (code completion, translation, summarization) but systematically fail at tasks requiring $R > 0.4$ (paradigm creation, contextual ethics, novel framework development). The 42% AI project abandonment rate occurs precisely when organizations attempt to deploy low-R systems for high-R challenges.

Architectural Isomorphism: Transformer architectures literally implement vector processing through tokenization (fragmentation), embedding (standardization), and attention (selective focus). They are mathematically constrained to $R < 0.3$ by their architectural foundations.

Measurement Convergence: Five independent approaches yield consistent R classifications, suggesting R captures a fundamental property rather than arbitrary categorization.

The choice facing individuals and institutions is thus genuinely binary: invest the energy required for cognitive sovereignty or accept thermodynamic dissolution into extractive systems. Physics doesn't negotiate.

4 Civilizational Thermodynamics: Why Spheres Require Surplus

The preceding sections established that cognitive sovereignty requires sustained energy investment exceeding entropy rate. But this framing, while thermodynamically precise, obscures a critical question: *where does the energy come from?* Individual cognition operates on approximately 21 watts—the brain's metabolic budget regardless of whether one is a Renaissance polymath or a micro-credential collector. The difference lies not in biological capacity but in *civilizational architecture*: how societies capture, concentrate, and allocate surplus energy toward cognitive development.

This section reveals what conventional analyses of education and expertise systematically ignore: the **hidden energy budget** that enabled historical sphere-building. Understanding this budget explains both why ancient and medieval systems produced multidimensional thinkers and why modern systems, despite unprecedented energy abundance, fail to do so.

4.1 The Hidden Energy Budget: What Hour-Counting Misses

Standard comparisons of educational systems count hours: 20 years of Greek philosophical training versus 3 years of Bologna bachelor’s degrees versus 40 hours of micro-credentials. These calculations, while accurate, miss the thermodynamic reality. *Not all hours are created equal.* The energy available for cognitive development depends on three factors that historical analyses systematically ignore.

4.1.1 Externalized Survival Costs

The medieval scholar at Oxford or Paris did not cook meals, maintain shelter, regulate body temperature, or navigate economic survival. These metabolic demands were *externalized*—absorbed by servants, college kitchens, heated common rooms, and institutional provisioning. The cognitive energy that a modern student expends on commuting, part-time employment, meal preparation, and domestic maintenance was, for historical scholars, available for synthesis and reflection.

Consider the thermodynamic accounting. A contemporary student’s waking hours might divide as:

- 6 hours: formal study (active cognition)
- 2 hours: commuting (metabolic maintenance, cognitive noise)
- 4 hours: employment (energy extraction by others)
- 3 hours: domestic labor (survival maintenance)
- 1 hour: fragmented leisure (insufficient for synthesis)

The medieval student’s allocation differed fundamentally:

- 6 hours: formal study (active cognition)
- 10 hours: “idle” time (synthesis, reflection, integration)
- 0 hours: employment, commuting, domestic labor (externalized)

The formal study hours appear identical. The *synthesis hours* differ by an order of magnitude.

4.1.2 Cognitive Idleness as Productive State

Jamadar’s (2025) finding that goal-directed cognition requires only 5% above baseline brain metabolism reveals something profound: the brain’s 20-watt baseline isn’t idle—it’s *processing*. The 95% of cognitive energy consumed during apparent rest performs

integration, pattern recognition, and cross-domain synthesis that cannot occur during focused attention.

This inverts conventional productivity thinking. The medieval student “wasting” 10 hours in contemplation, conversation, and wandering was not failing to study but allowing baseline processing to integrate what focused study had captured. The 5% sprint of active cognition *refines and formalizes* what the 95% baseline *generates*. Modern educational systems, optimizing for measurable active learning, systematically eliminate the generative baseline while preserving only the refinement sprint.

Wiehler et al. (2022) provide the biochemical mechanism: sustained cognitive control causes glutamate accumulation in the lateral prefrontal cortex. Mental fatigue is not metaphor but metabolic reality. The brain requires recovery periods not from laziness but from chemistry. Historical educational systems, through apparent inefficiency, provided this recovery. Modern systems, through apparent efficiency, prevent it.

4.1.3 Environmental Energy Subsidies

Thermoregulation consumes significant metabolic energy. A body maintaining 37°C in cold environments diverts calories from cognition to heat generation. Historical centers of learning cluster in climatically favorable regions (Mediterranean academies, mild English college towns) or provided environmental energy subsidies (hypocaust heating in Roman schools, great hall fires in medieval colleges).

The energy equation:

$$E_{cognitive} = E_{total} - E_{survival} - E_{thermoregulation} - E_{extraction} \quad (2)$$

Historical systems minimized the subtractive terms for scholars through institutional architecture. Modern systems, despite central heating and abundant food, maximize extraction ($E_{extraction}$) through employment demands, leaving cognitive surplus at or near zero for most learners.

4.2 Energy Transitions and Cognitive Architecture

The history of human cognition cannot be separated from the history of energy capture. Smil’s (2017) framework—viewing civilization as “the quest for controlling greater stores and flows of more concentrated and versatile forms of energy”—applies directly to cognitive development. Each major energy transition reshaped the possibilities for sphere-building.

4.2.1 The Agricultural Foundation (10,000 BCE–500 CE)

The Neolithic Revolution enabled the first cognitive specialists. Hunter-gatherer bands, capturing approximately 10,000 kcal/day per capita through foraging, could not sustain non-productive members. Agricultural societies, capturing 20,000–40,000 kcal/day per capita, generated surplus sufficient to release priests, scribes, and philosophers from direct food production.

Greek philosophical schools emerged from this surplus. The *paideia* of 20-year comprehensive education required that students neither farm nor fight. Athenian democracy’s leisure class—enabled by slave labor capturing solar energy through agriculture—provided the metabolic foundation for Plato’s Academy and Aristotle’s Lyceum. The energy was somatic (human and animal muscle), the surplus modest, but the *concentration* sufficient for sphere cultivation.

4.2.2 The Medieval Energy Transition (500–1500 CE)

The collapse of Roman centralized extraction initially devastated cognitive institutions. But the medieval period witnessed a critical energy transition: widespread adoption of water and wind power.

The Domesday Book of 1086 records over 5,000 watermills in England alone—one of history’s largest capital investments in non-somatic energy capture. Cistercian monasteries pioneered watermill technology, becoming what historians recognize as Europe’s first transnational corporations. This extrasomatic energy freed human labor from grinding, fulling, and sawing, creating surplus for other purposes.

Critically, medieval institutions *concentrated* this surplus rather than distributing it. The Church—owning an estimated 20–33% of all cultivated land in Western Europe—functioned as the primary mechanism for energy extraction and allocation. The tithe system collected 10% of agricultural production not in currency but in *calories*: grain, livestock, eggs. This was energy taxation at civilizational scale.

Where did this extracted energy go? Partially into “frozen” form—the Gothic cathedrals that consumed an estimated 21.5% of the Paris Basin’s GDP between 1100 and 1250 (Siebert, 2019). But also into *cognitive institutions*: the monasteries that preserved texts, the cathedral schools that taught the trivium, the universities that emerged from 1088 onward.

The correlation between energy transition and university founding is not coincidental:

- Watermill proliferation: 1000–1200 CE
- University founding wave: Bologna (1088), Paris (c.1150), Oxford (c.1167)
- Medieval Warm Period peak: 1000–1300 CE

Agricultural surplus, extrasomatic energy capture, favorable climate, and concentrated institutional allocation converged to enable comprehensive education lasting 15 years (7 years trivium/quadrivium + 8 years specialized faculty).

4.2.3 The Climate Correlation

Climate operates on cognitive production not directly but through the mediation of agricultural surplus. The Medieval Warm Period (900–1300 CE), with temperatures 1.0–1.4°C above twentieth-century baselines, extended growing seasons, enabled northern cultivation, and approximately doubled European population between 1000 and 1300.

This demographic-agricultural surplus funded the first university wave. The correlation is precise: major university foundations cluster in the MWP’s favorable window.

The Little Ice Age (1300–1850 CE) presents an apparent paradox. Climate deterioration caused the Great Famine (1315–17), amplified Black Death mortality (1347–52), and disrupted agricultural production. Yet the Scientific Revolution—Copernicus, Galileo, Kepler, Descartes, Newton—occurred during the LIA’s coldest phases.

The resolution lies in **institutional persistence** and **surplus concentration**. Once universities, printing presses, and scholarly networks existed, intellectual production continued through adverse conditions via geographic substitution, lagged output, and institutional resilience. Moreover, the LIA’s crises *concentrated* resources among survivors and intensified state competition for knowledge advantages.

The climate-cognition relationship follows a threshold model rather than linear correlation:

Below a certain surplus threshold, societies cannot afford large scholar populations. Once above threshold, other factors dominate: institutional allocation, state competition, technological infrastructure.

4.3 Christianity as Thermodynamic Engine

The Church’s role in medieval cognitive development extends beyond institutional patronage. Christianity functioned as a **thermodynamic engine**—a system for extracting energy from the general population and concentrating it in specialized institutions.

4.3.1 The Extraction Mechanism

Medieval Christianity deployed what we might call “ideological efficiency enhancement.” The doctrine of salvation as delayed gratification—infinite utility in the afterlife exchanged for current labor—effectively lowered the “maintenance cost” of the workforce. Peasants accepted subsistence consumption not from coercion alone but from belief that present suffering represented investment in future reward.

Irenaeus articulated the mechanism explicitly: “It is just that in that very creation in which they toiled and were afflicted... they should receive the reward of their suffering... and that in the very creation in which they endured servitude, they should also reign.” This theology framed labor as *investment* rather than exploitation, enabling extraction rates impossible through violence alone.

The result: surplus that would otherwise have been consumed (peasant demand for better conditions) or dissipated (inefficient coercion costs) became available for institutional allocation.

4.3.2 The Allocation Mechanism

Extracted surplus flowed through ecclesiastical channels to specific destinations:

- **Monastic preservation:** Scriptoria maintaining textual transmission across centuries
- **Cathedral schools:** Educating clergy who would staff administrative and intellectual positions
- **Universities:** From 1088, formal institutions for comprehensive knowledge development
- **Frozen energy:** Cathedrals as permanent displays of extraction capacity, reinforcing the ideology enabling extraction

The feedback loop was self-sustaining: theology justified extraction; extraction funded institutions; institutions produced theologians; theologians refined extraction-justifying theology.

4.3.3 The Cognitive Consequence

This system—for all its ethical problems from modern perspectives—*produced spheres*. The medieval scholar, supported by the surplus labor of hundreds or thousands of peasants concentrated through ecclesiastical taxation, could spend 15 years building multidimensional cognitive architecture. The energy came from agriculture; the concentration came from ideology; the allocation came from institutional design.

Modern systems lack not energy but *concentration mechanisms*. Total energy availability vastly exceeds medieval levels. But this energy distributes across populations rather than concentrating in cognitive institutions. The medieval peasant’s surplus funded monasteries; the modern worker’s surplus funds consumption. No contemporary ideology legitimizes the extraction and concentration that sphere-building requires.

4.4 The Polymath Production Curve

Burke’s (2020) cultural history of polymathy provides empirical validation of the thermodynamic framework. Polymath production follows a bell curve peaking in the seventeenth century—precisely when we would expect given the energy-surplus-concentration model.

4.4.1 The Timeline

- **1350–1600 (Renaissance Rise):** The *uomo universale* ideal emerges. Artist-engineers like Leonardo embody cross-domain mastery. Energy remains primarily somatic, but post-plague labor surplus allows urban elites to fund individual genius through patronage.
- **1600–1700 (Peak Production):** Burke identifies this as the “Age of Erudition,” producing “monsters” like Leibniz, Athanasius Kircher, and Newton. The Republic of Letters—a transnational scholarly network—flourishes under royal and aristocratic patronage. Concentrated surplus from state centralization and colonial extraction funds comprehensive intellectual development.
- **1700–1850 (Tension and Decline):** Knowledge volume strains individual capacity. “Civil servant polymaths” like Goethe and von Humboldt use state positions to fund diverse interests. But the strain is evident; the sphere becomes harder to maintain.
- **1850–Present (Collapse):** The word “scientist” is coined in 1833, marking disciplinary fragmentation. Burke notes he “cannot identify any [polymaths] born after 1960.” Specialization becomes not merely common but *mandatory*.

4.4.2 The Economic Mechanism

Polymath production correlates with **concentrated surplus under patronage**, not with general prosperity or distributed wealth. The seventeenth-century “monsters of erudition” emerged during the Little Ice Age—not despite but *because of* its crises.

Consider Athanasius Kircher, who wrote forty massive volumes spanning geology, music theory, Egyptology, and magnetism. Kircher’s position at the Roman College required the surplus labor of thousands of peasants, extracted through Church taxation and concentrated in a single institution. He was, in energy terms, a *luxury good*—sustainable only through radical inequality of allocation.

The critical distinction: **patronage rewarded breadth** (the glory of the patron reflected in the universality of their client); **markets reward depth** (efficiency demands specialization). As economies shifted from feudal/mercantile to industrial/capitalist, the

economic incentive for polymathy vanished. No market participant will pay for comprehensive knowledge when specialized knowledge costs less and produces immediate returns.

4.4.3 The Ouroboros Pattern

The polymaths consumed themselves. This is the civilizational tragedy our framework reveals:

1. Seventeenth-century polymaths, funded by concentrated surplus, developed comprehensive knowledge
2. Their work enabled the Scientific Revolution
3. The Scientific Revolution enabled the Industrial Revolution
4. The Industrial Revolution massively increased energy capture (coal, steam)
5. Increased energy capture exponentially increased information production
6. Information explosion exceeded individual cognitive capacity
7. Specialization emerged as efficiency response to cognitive overload
8. Specialists replaced polymaths
9. Specialized vectors became extractable by algorithms
10. AI now harvests the patterns that polymaths inadvertently made possible

The polymaths ate their own tail. They created the conditions that made polymathy impossible. They lit the fire that consumed them.

Burke captures the paradox: “The foundation of new disciplines in an age of specialization offered a new role for polymaths, at least in the short term, since a new discipline is necessarily taught in the first generation by professors who have been trained in something else.” Polymaths are the *spark* of new systems; specialists are the *maintenance*. But maintenance workers systematically eliminate the conditions that produced sparks.

4.5 The Modern Energy Paradox

Contemporary civilization commands energy resources that would astound medieval observers. Per capita energy consumption in developed nations exceeds 100 kWh/day—four times medieval European levels. Yet cognitive surplus for sphere-building approaches zero. The paradox demands explanation.

4.5.1 Distribution vs. Concentration

Medieval systems concentrated modest surplus in cognitive institutions. Modern systems distribute abundant energy across populations. The Cistercian monastery captured thousands of peasants' surplus and directed it toward a single center of learning and preservation. The modern economy distributes energy as wages, which workers then allocate to consumption, housing, transportation, and entertainment.

No contemporary mechanism concentrates surplus for cognitive development at civilizational scale. Universities, rather than receiving concentrated societal surplus, compete in markets for tuition revenue and research grants. Their incentive is efficiency (minimize costs, maximize throughput) rather than cultivation (maximize cognitive development regardless of cost).

4.5.2 Extraction Without Investment

The platform economy represents extraction without corresponding investment. Zuboff's (2019) "surveillance capitalism" harvests behavioral data—the organized patterns of human cognition—without investing in the cognitive development that created those patterns. This is thermodynamic parasitism: consuming negentropy while contributing only entropy.

The medieval Church, for all its extractive capacity, *reinvested* in cognitive institutions. Tithes funded monasteries that preserved texts and educated clergy. Modern platforms extract without reinvestment. The cognitive surplus they harvest depletes; they make no provision for replenishment.

4.5.3 The Survival Trap

Most critically, modern workers must allocate their cognitive energy to survival rather than synthesis. The hidden energy budget that enabled historical sphere-building—externalized survival costs, abundant synthesis time, environmental subsidies—has been systematically eliminated.

The contemporary knowledge worker:

- Commutes (cognitive energy dissipated in noise and stress)
- Works for others (cognitive energy extracted by employers)
- Maintains domestic life (cognitive energy consumed by survival)
- Experiences fragmented leisure (insufficient contiguous time for synthesis)
- Has **zero hours** available for the integration that baseline processing requires

The medieval scholar, the Renaissance patron’s client, and even the nineteenth-century “civil servant polymath” had what contemporary workers lack: *protected time for cognitive integration*. This is not luxury but thermodynamic necessity. Without it, the 95% baseline processing that generates insight has nothing to work with; the 5% active cognition has nothing to refine.

4.6 Threshold Analysis: The Approaching Collapse

The thermodynamic framework enables prediction. If cognitive sovereignty requires energy investment exceeding entropy rate, and if modern systems systematically prevent such investment, then collapse becomes calculable.

4.6.1 The Critical Threshold

Tainter’s (1988) analysis of civilizational collapse identifies a critical threshold: when the marginal return on complexity investment becomes negative, societies simplify—often catastrophically. Our framework translates this to cognitive terms: when the population’s average resistance to extraction ($\bar{R}$) falls below approximately 0.05, complex civilization becomes unsustainable.

Current trajectories suggest this threshold approaches within decades rather than centuries. The decay rate in $\bar{R}$ has accelerated from 0.002/year (1800–1950) to 0.0054/year (1950–2024). Extrapolation is hazardous, but the direction is unambiguous.

4.6.2 The Reconstruction Window

The cohort that experienced pre-Bologna education—those who built comprehensive foundations before specialization—ages out by 2040–2050. They carry embodied knowledge of how spheres were built: not merely information about curricula but the lived experience of intellectual integration across domains.

Without deliberate preservation and transmission, this knowledge becomes archaeological. We would face not reform but reconstruction from fragments, attempting to reverse-engineer what we deliberately destroyed.

4.6.3 The Binary Choice

The thermodynamic analysis reveals no middle path. Either:

1. **Concentrated reinvestment:** Societies create mechanisms to concentrate surplus in cognitive development, accepting the inequality this implies, rebuilding sphere-cultivating institutions

2. **Distributed dissolution:** Societies continue distributing energy across consumption, allowing cognitive architecture to degrade to algorithmic substitutability, accepting replacement by systems that process vectors more efficiently than degraded humans

The German engineering faculties maintaining five-year integrated programs despite Bologna pressure, the few remaining comprehensive curricula, the individual practitioners building “local negentropy bubbles”—these represent choice (1) at micro scale. Whether macro-scale reconstruction remains possible is the question this analysis cannot answer but must pose.

4.7 Summary: The Civilizational Thermodynamics of Cognition

The sphere-to-vector transformation documented in preceding sections reflects civilizational energy dynamics, not merely educational policy choices. Historical sphere-building required:

1. **Energy surplus:** Agricultural and extrasomatic energy capture exceeding subsistence needs
2. **Concentration mechanisms:** Institutions (Church, patronage, state) directing surplus toward cognitive development rather than distributed consumption
3. **Hidden energy budget:** Externalized survival costs, protected synthesis time, environmental subsidies enabling the 95% baseline processing that generates insight
4. **Temporal investment:** Decades of comprehensive development before specialization

Modern civilization possesses (1) in abundance but has eliminated (2), (3), and (4). The result: maximum energy availability producing minimum cognitive surplus. We have industrialized extraction while eliminating investment. We have optimized efficiency while destroying the inefficiency that synthesis requires. We have more energy than any previous civilization and less capacity for the thinking that matters.

The polymaths who built our world consumed themselves in building it. The Scientific Revolution they enabled created the Industrial Revolution that generated the information explosion that made polymathy impossible. Now their successors—specialists trained in vectors, documented in procedures, optimizable by algorithms—face replacement by systems that process patterns more efficiently than biological cognition ever could.

Physics doesn’t negotiate. Neither does civilizational thermodynamics. The question is not whether the analysis is correct but whether, knowing it, we can choose differently than the trajectory our systems have inscribed.

5 Historical Analysis: From Cognitive Cathedrals to Vector Factories

5.1 The Inherited Wholeness (500 BCE - 1829)

The historical record reveals a stunning truth: for over two millennia, human intellectual development operated on fundamentally different thermodynamic principles than today’s educational systems. The evidence spans from Aristotle’s Lyceum to the death of Thomas Young in 1829, marking what [Robinson \(2006\)](#) definitively identifies as “The Last Man Who Knew Everything.”

5.1.1 Ancient Foundations: The 20-Year Investment

The Greek philosophical schools established a pattern that would persist for two millennia: knowledge required decades of energy investment. Aristotle’s students at the Lyceum underwent 20 years of comprehensive education before specialization. As [Mouzala \(2024\)](#) demonstrate in their interdisciplinary analysis, understanding Greek cognitive categories now requires seventeen contemporary specialists: what individual ancient scholars grasped intuitively. [Erkizan \(1997\)](#) dissertation on nous alone spans hundreds of pages analyzing a single term that Aristotle’s students understood through lived practice.

This wasn’t primitive education; it was high-energy cognitive architecture. Students developed not four categories of knowing (our contemporary DIKW pyramid) but over ten distinct types: episteme (theoretical knowledge), techne (craft knowledge), phronesis (practical wisdom), metis (cunning intelligence), nous (intellectual intuition), sophia (theoretical wisdom), gnosis (spiritual knowledge), synesis (comprehension), episteme praktike (practical science), and dianoia (discursive reasoning).

5.1.2 Medieval Synthesis: The Cathedral Builders

Medieval universities, beginning with Bologna (1088), Oxford (1167), and Paris (1150), institutionalized cognitive wholeness through mandatory comprehensive foundations before permitting any specialization. The system exemplified sphere-first architecture: seven years building multidimensional connectivity, then eight years deepening domain expertise on that foundation.

The Trivium (Years 1-4):

- Grammar: Deep linguistic architecture in Latin, Greek, often Hebrew
- Logic: Universal analytical capability across all domains
- Rhetoric: Synthesis and persuasion, weaving knowledge into compelling narrative

The Quadrivium (Years 4-7):

- Arithmetic: Number theory and divine proportion
- Geometry: Spatial reasoning through memorized Euclid
- Music: Mathematical harmony as physics and theology
- Astronomy: Navigation, cosmic cycles, understanding place in universe

This seven-year foundation was *prerequisite* and *non-negotiable*. University statutes explicitly prohibited advancement to specialized faculties without demonstrated mastery. As the University of Paris declared: “No student shall be admitted to the study of theology who has not first completed the full course in the seven liberal arts.” Only after receiving the *Magister Artium* (Master of Arts) degree—documenting comprehensive foundation—could students spend an additional 8-12 years specializing in theology, law, or medicine.

The complete medieval educational system thus invested approximately 15 years total: 7 years (47%) building sphere architecture, 8 years (53%) extending vectors from that foundation. This produced specialists who retained adaptive capacity—theologians who understood mathematics, physicians who grasped philosophy, lawyers versed in logic and rhetoric.

Medieval guilds paralleled this architecture. As [De la Croix et al. \(2018\)](#) document, guild apprenticeships achieved “knowledge transmission that transcended kinship boundaries,” creating distributed cognitive networks maintaining both depth and diversity. A master carpenter’s 7-10 year apprenticeship included comprehensive material science, applied geometry, aesthetic theory, and economic reasoning before specialization in furniture, ship-building, or architectural carpentry.

The medieval synthesis maintained high foundation-to-specialization ratios (40-50%), enabling what we now recognize as cognitive sovereignty: specialists grounded in spherical architectures, resistant to extraction precisely because their expertise emerged from multidimensional foundations rather than narrow optimization.

5.1.3 The Last Universal Minds

Thomas Young (1773-1829) represents the definitive terminus of classical polymathy. When invited to write for the Encyclopedia Britannica, Young offered expertise in: “Alphabet, Annuities, Attraction, Capillary Action, Cohesion, Colour, Dew, Egypt, Eye, Focus, Friction, Halo, Hieroglyphic, Hydraulics, Motion, Resistance, Ship, Sound, Strength, Tides, Waves, and anything of a medical nature” ([Robinson, 2006](#)). His death marks the last moment when comprehensive mastery across human knowledge remained possible.

5.2 The Great Amputation (1750-1920)

The destruction of cognitive wholeness occurred through identifiable phases, each characterized by exponentially decreasing energy investment in knowledge formation.

5.2.1 The Architects of Reduction

Three figures provided the intellectual blueprints for cognitive standardization:

Adam Smith (1776): The pin factory model didn't just divide labor; it divided cognition. Eighteen steps to make a pin became the template for fragmenting any complex process. The pin-maker who once understood metallurgy, aesthetics, and markets became an operative who knew only "step seven: straightening wire."

Frederick Taylor (1911): Arrived not as destroyer but as optimizer of existing wreckage. Workers had already lost craft knowledge to 135 years of industrialization. Taylor measured the poverty and called it science. His death in 1915 (nearly broke despite his efficiency expertise) suggests even he couldn't optimize what had been destroyed.

Russell Ackoff (1989): The DIKW pyramid reduced millennia of cognitive diversity to four categories: Data→Information→Knowledge→Wisdom. Intended as helpful simplification, it became the extraction template for all subsequent knowledge management systems.

5.2.2 The Polymath Extinction Event

[Burke \(2020\)](#) comprehensive analysis provides the authoritative timeline:

- 17th century: "Golden age" of polymathy
- 18th century: Progressive specialization begins
- 19th century: University disciplines formalize boundaries
- 1960: Last polymaths born (Burke: "I am unable to identify any [polymaths] who were born after the year 1960")

The quantum generation (1925-1927) represents the final moment when individuals could create paradigm shifts. As [Beller \(1996\)](#) documents, after Einstein, Bohr, Heisenberg, and Schrödinger, physics required teams and apparatus. Von Neumann (d. 1957), Russell (d. 1970), and Polanyi (d. 1976) were "the last intellectual giants [who] kept polymathy's veins flowing with blood until they themselves finally flatlined, and it did too" ([Hoel, 2025](#)).

5.3 The Bologna Massacre (1999-2024)

The Bologna Declaration of June 19, 1999, represents the industrialization of education at continental scale. Signed by 29 European education ministers, it promised “harmonization.” The documentation reveals systematic cognitive destruction.

5.3.1 The Standardization Mechanism

Bologna’s tools of cognitive standardization:

- **Modularization:** Knowledge fractured into discrete, assessable units
- **ECTS Credits:** Learning literally becomes currency (48 million credits traded annually)
- **Learning Outcomes:** Every thought must be measurable and predetermined
- **Competency Frameworks:** Humans described as standardized skill-containers

As [Gleeson \(2021\)](#) documents, ECTS evolved from mobility tool to “market commodity,” with education value shifting from mastery to “quantifiable outputs.”

5.3.2 The Destruction Metrics

The empirical evidence of degradation:

System Thinking Loss: [Kaiser and Schröder \(2022\)](#) confirm that “immediate factors such as Systems Thinking, collaboration and communication...are not explicitly addressed, although they are considered essential” in post-Bologna engineering curricula.

Industry Testimony:

- “We sacrificed the Diplom-Ingenieur with heavy hearts for a greater goal, namely international connectivity” (VDI President Ungeheuer, 2016)
- TU Dresden Faculty: “The university Bachelor’s degree in six semesters does not lead to a professionally qualifying degree” ([Odenbach and Krauthäuser, 2015](#))
- German collective bargaining: Different wage groups for Diplom vs Bachelor/Master holders ([Wieschke et al., 2020](#))

Cognitive Diversity Collapse:

- Medieval graduate: 15+ cognitive categories active
- Pre-Bologna graduate: 8-10 categories maintained
- Post-Bologna graduate: 3-4 categories (all episteme variants)

- Diversity loss: 73-80%

Social Stratification: [Kroher et al. \(2021\)](#) demonstrate Bologna “generated new forms of social inequalities,” with lower-background students attending Masters programs less frequently: the opposite of democratization claims.

5.4 The Contemporary Harvest (2020-2024)

The current AI revolution represents not disruption but collection: harvesting the pre-vectorized knowledge that Bologna prepared.

5.4.1 The Implementation Crisis

[S&P Global Market Intelligence \(2025\)](#) provides the smoking gun: 42% of companies abandoned most AI initiatives in 2024, up from 17% the previous year. Organizations scrapped an average of 46% of proof-of-concepts before production. This isn’t technological failure; it’s the discovery that vectorized knowledge lacks the contextual depth AI supposedly replaces.

[McKinsey & Company \(2025\)](#) confirms only 1% of companies consider themselves “mature” in AI deployment, with fewer than 10% of use cases progressing past pilot stage. The “pilot purgatory” reveals a fundamental mismatch: AI trained on standardized patterns cannot navigate the complex domains where human judgment remains essential.

5.4.2 The Extraction Confessions

Industry now admits what critics predicted:

IBM’s Reversal: After laying off 8,000 HR employees for AI replacement, CEO Krishna revealed: “Our total employment has actually gone up, because what [AI] does is it gives you more investment to put into other areas” ([Krishna, 2024](#)). The 6% of tasks requiring “empathy, nuance, trust” proved unextractable.

Duolingo’s Collapse: CEO von Ahn: “We’d rather move with urgency and take occasional small hits on quality than move slowly” ([von Ahn, 2024](#)). Users with 1,131-day streaks canceled in protest. Content became “repetitive, robotic” without the sphere-holders who created pedagogical soundness.

Microsoft’s Quantification: Chief Commercial Officer Althoff celebrated “\$500 million saved using AI in 2024” ([Althoff, 2024](#)), explicitly commodifying human knowledge as extractable value while conducting 15,000 layoffs.

5.5 The Thermodynamic Gradient

The historical timeline reveals exponential decay in cognitive power investment:

Era	Formation Duration	Power×Time (W·years)	R	C_S (relative)
Pre-Industrial (pre-1800)	15-20 years	30-40	0.7-0.8	1.00 (baseline)
Industrial (1800-1950)	5-7 years	10-14	0.4-0.5	0.33
Post-War (1950-1999)	4-5 years	8-10	0.2-0.3	0.17
Bologna Era (1999-2020)	3 years	6	0.1-0.2	0.10
Micro-credentials (2020+)	hours-days	< 0.1	< 0.1	< 0.01

Table 2: Historical decline in cognitive sovereignty components. R (Resistance to Extraction) quantifies architectural complexity; C_S (Cognitive Sovereignty) normalized to pre-industrial baseline. Energy investment assumes 2W average cognitive work rate.

This trajectory isn’t evolution—it’s entropy acceleration. Each phase reduced energy invested while claiming improved “efficiency,” approaching Tainter’s collapse threshold ($R < 0.05$) where complex civilization becomes thermodynamically unsustainable.

5.6 The Pattern Crystallizes

The historical analysis reveals a consistent template:

1. **Helpful Framework:** Simplification for management (pin factory, DIKW, Bologna)
2. **Institutional Adoption:** Scale across systems (industrialization, universities)
3. **Standardization:** Eliminate variation (taylorism, ECTS)
4. **Optimization:** Perfect the poverty (metrics, rankings)
5. **Extraction:** Harvest the patterns (AI training)

Each step appeared rational. Together, they constitute systematic cognitive dismemberment. The guild master would recognize this pattern: it’s exactly how craft knowledge died. First documentation, then systematization, then optimization, then obsolescence.

But history also reveals what resists extraction: the phronesis of contextual judgment, the metis of adaptive cunning, the nous of intuitive leaps. These require energy investment that cannot be modularized, standardized, or extracted. They exist only in the sustained practice of cognitive sovereignty, precisely what our educational systems no longer provide.

The timeline is unforgiving: 2,500 years building cognitive architecture, 250 years dismantling it, 25 years completing the destruction. The feast has begun, but we prepared the meal ourselves.

6 Contemporary Evidence: The Thermodynamic Collapse in Real-Time

6.1 The Implementation Crisis: Entropy Acceleration

The transition from theory to implementation reveals thermodynamic reality: systems optimized for extraction cannot sustain themselves. [S&P Global Market Intelligence \(2025\)](#) provides the quantitative evidence: 42% of companies abandoned the majority of their AI initiatives in 2024, a surge from 17% the previous year. Organizations scrapped an average of 46% of proof-of-concepts before reaching production.

This isn't technological failure; it's entropy manifestation. Systems trained on vectorized knowledge lack the energetic foundation to navigate complex domains. [McKinsey & Company \(2025\)](#) confirms that fewer than 10% of deployed use cases progress beyond pilot stage, with only 1% of companies achieving "mature" AI deployment. The "pilot purgatory" represents thermodynamic reality: extractive systems consuming their own foundations.

The failure pattern follows predictable entropy acceleration:

- Initial enthusiasm (energy investment appears minimal)
- Pilot success (controlled conditions mask entropy)
- Scaling attempts (complexity emerges, vectors fail)
- Abandonment (entropy overwhelms system)
- Denial and repetition (new initiatives, same pattern)

6.2 The Extraction Disasters: Consuming the Foundations

6.2.1 The Duolingo Cliff Fall

In April 2024, Duolingo demonstrated Cynefin's cliff in action. The company eliminated over 100 contract writers, translators, and curriculum experts, the sphere-holders who created pedagogically sound content. Their replacement: OpenAI's GPT models trained on the extracted patterns.

The thermodynamic collapse was immediate:

- Users with 1,131-day streaks canceled in protest
- 6.7 million TikTok followers witnessed brand suicide
- Content quality degraded from pedagogical design to pattern matching

- The company deleted social media accounts rather than face backlash

This represents more than business failure. Language learning requires *phronesis* (contextual judgment), *metis* (cultural navigation), and *nous* (intuitive understanding), none extractable through pattern analysis. The vectors could replicate surface grammar but not the sphere of cultural embodiment that makes language acquisition possible.

6.2.2 IBM’s Forced Rehiring: The Confession

IBM’s experiment provided controlled evidence of extraction limits. After laying off approximately 8,000 HR employees for AI replacement, the company was forced to rehire human workers. The AI systems could execute procedures but couldn’t navigate the organizational complexity requiring genuine human judgment.

This validates [Collins \(2010\)](#) taxonomy of tacit knowledge:

- **Relational Tacit Knowledge:** Extractable (procedures, rules)
- **Somatic Tacit Knowledge:** Partially extractable (physical skills)
- **Collective Tacit Knowledge:** Unextractable (social embedding required)

IBM discovered that HR work is primarily CTK: requiring authentic participation in organizational culture. No amount of data extraction could replicate the energetic investment of lived organizational experience.

6.3 The Cognitive Degradation: Entropy in Human Systems

[Rinta-Kahila et al. \(2023\)](#) documented the entropy mechanism in accounting firms: “Cognitive automation leads to complacency and reduces mindfulness in tasks, gradually eroding essential skills.” The degradation follows thermodynamic principles: systems not actively maintained through energy investment inevitably decay.

The vicious cycle accelerates:

1. Automation reduces human practice (energy disinvestment)
2. Reduced practice erodes skills (entropy increases)
3. Eroded skills increase dependency (system brittleness)
4. Dependency accelerates automation (further disinvestment)
5. System becomes irreversibly degraded (thermodynamic collapse)

Contemporary evidence reveals this pattern across industries:

- **Radiology:** AI-assisted diagnosis reducing pattern recognition capabilities

- **Legal:** Document automation eliminating analytical skill development
- **Finance:** Algorithmic trading destroying market intuition
- **Education:** Automated grading eliminating pedagogical judgment

Each represents entropy acceleration through energy disinvestment. The skills don't transfer to machines; they simply cease to exist.

6.4 The Quantification of Extraction

Microsoft's Chief Commercial Officer Judson Althoff provided explicit commodification: "\$500 million saved using AI in 2024, and that's just at its call centers" ([Althoff, 2024](#)). This represents direct quantification of extracted human cognitive value, with 15,000 layoffs converting knowledge workers into cost reductions.

The thermodynamic calculation:

- Human cognitive development: 10-20 years energy investment
- Knowledge extraction period: 6-12 months
- AI replication quality: 60-70% of original
- Entropy acceleration: Exponential
- System sustainability: Approaching zero

Shopify CEO Tobi Lütke institutionalized the entropy through policy: employees must prove why they "cannot get what they want done using AI" before requesting headcount. The burden of proof inverted: humans must justify their existence against systems that demonstrably fail at 42% abandonment rates.

6.5 The Domain Confusion Catastrophe

The Cynefin framework diagnostic reveals systematic domain confusion driving contemporary failures:

Clear Domain (Best Practices)

- Where vectors work: Repetitive, rule-based tasks
- AI performance: Adequate
- Human advantage: None

Complicated Domain (Good Practices)

- Where expertise matters: Technical analysis, diagnostics

- AI performance: Limited success with sufficient data
- Human advantage: Contextual judgment

Complex Domain (Emergent Practices)

- Where spheres essential: Innovation, relationship management, crisis navigation
- AI performance: Systematic failure
- Human advantage: Irreplaceable

Chaotic Domain (Novel Practices)

- Where only humans function: Crisis response, paradigm creation
- AI performance: Complete failure
- Human advantage: Absolute

Organizations systematically misclassify complex work as complicated, attempting to solve sphere problems with vector solutions. As [Kempermann \(2017\)](#) warns: “Complex problems in biomedicine are often treated as if they were actually not more than the complicated sum of solvable sub-problems. . . [with] dangerous consequences, especially in clinical contexts.”

6.6 The Micro-Credential Delusion

[Lumina Foundation \(2025\)](#) celebrates: 96% of employers say micro-credentials strengthen job applications, 90% offer 10-15% higher starting salaries. Yet this represents preference for granular verification over meaningless degrees, not actual capability development.

[Ha et al. \(2022\)](#) systematic review reveals the truth: while 13 of 14 studies show positive outcomes, these measure satisfaction not competence. [Joshi \(2019\)](#) provides empirical evidence: bootcamps help non-technical graduates enter tech but provide “minimal benefit for those already holding technical degrees”; you can’t accelerate what doesn’t exist.

The thermodynamic reality:

- Google Career Certificate: 3-6 months, single skill
- Coursera Specialization: 4-8 weeks, narrow domain
- LinkedIn Learning Path: 5-10 hours, procedural knowledge
- Energy investment: Approaching zero
- Actual expertise developed: Statistical noise

6.7 The Pattern Crystallizes

Contemporary evidence reveals not technological limitation but thermodynamic law: systems that extract without investing inevitably collapse. Every failure represents entropy overwhelming extractive ambition. Every abandonment confirms that spheres cannot be sustained through vectors alone.

The feast isn't succeeding; it's creating mutual destruction:

- Organizations destroying human capabilities for non-functional AI
- Workers losing skills whether or not replacements function
- Systems becoming simultaneously de-skilled and non-automated
- Entropy accelerating through positive feedback loops

We're witnessing not technological revolution but thermodynamic collapse: the inevitable consequence of attempting to sustain complex systems through extraction rather than investment.

7 The AI Mirror: How We Built Machines in Our Own Vectorized Image

7.1 The Architecture Isomorphism

Large Language Models didn't emerge from technological innovation but from educational standardization. Their architecture (tokenization, embedding, attention) represents the precise computational implementation of the cognitive patterns we've been training into humans since Bologna. The mirror is perfect because we designed both sides to match.

[Vaswani et al. \(2017\)](#) unknowingly documented this isomorphism in "Attention is All You Need." They described a transformer architecture that exactly parallels the educational transformations we've imposed on human cognition. Each component of the LLM corresponds to a deliberate modification of human thinking patterns implemented through systematic educational reform.

The correspondence isn't metaphorical; it's technical. We trained humans to process information in discrete, standardized, assessable units. Then we built machines that process information in discrete, standardized, assessable units. The machines work because we've spent decades preparing their training data: humans thinking in machine-compatible patterns.

7.2 Tokenization: The Modularization Mirror

7.2.1 Educational Tokenization (1999-Present)

The Bologna Process fragmented knowledge into European Credit Transfer System (ECTS) units: discrete, tradeable, stackable tokens of learning. A bachelor's degree became 180 tokens. A master's became 120 tokens. Knowledge literally became countable units divorced from integrated understanding.

Each credit represents 25-30 hours of “student workload”: not comprehension, not wisdom, not capability, but time units converted to knowledge tokens. Students collect tokens, institutions validate tokens, employers evaluate token counts. The system processes tokens, not understanding.

7.2.2 Computational Tokenization

LLMs fragment language into tokens—discrete units typically representing 4-5 characters or common word fragments. “Understanding” becomes “Under” + “stand” + “ing”: three tokens with no inherent meaning, only statistical relationships to other tokens.

The process is identical:

- **Input:** Continuous human thought/language
- **Transformation:** Fragmentation into discrete units
- **Processing:** Statistical manipulation of fragments
- **Output:** Reconstructed appearance of coherence

Both systems destroy wholeness to create processability. The medical student who once understood physiology as integrated system now processes cardiovascular (7.5 ECTS), respiratory (7.5 ECTS), and endocrine (7.5 ECTS) as separate tokens. The LLM that processes “heart” + “beat” as separate tokens mirrors the student processing organs as isolated credits.

7.3 Embedding: The Standardization Mirror

7.3.1 Educational Embedding (Competency Frameworks)

Post-Bologna education embeds diverse human capabilities into standardized competency vectors. The European Qualifications Framework defines eight levels across three dimensions (knowledge, skills, responsibility/autonomy), creating a 24-dimensional space where every human capability must find coordinates.

A master carpenter's embodied wisdom (decades of wood grain intuition, tool extension into consciousness, weather prediction through timber behavior) becomes:

- Knowledge: Level 5 (“comprehensive, specialized, factual and theoretical”)
- Skills: Level 5 (“comprehensive range of cognitive and practical skills”)
- Autonomy: Level 5 (“exercise management and supervision”)

The infinite-dimensional sphere of craft mastery compressed to a point in 24-dimensional competency space.

7.3.2 Computational Embedding

LLMs embed tokens into vector spaces, typically 768-1536 dimensions where each word/concept receives fixed coordinates. “Love” might be $[0.23, -0.45, 0.67, \dots]$, forever frozen at those coordinates regardless of context. Cleopatra’s love for Antony, a mother’s love for her child, and “I love pizza” all map to the same vector, distinguished only through attention mechanisms.

The parallel is exact:

- **Pre-standardization:** Infinite contextual meaning
- **Embedding process:** Forced mapping to fixed coordinates
- **Result:** Standardized vectors that can be computed but lose essence

Both systems convert qualitative richness into quantitative poverty. The embedding makes computation possible by destroying exactly what made the original valuable.

7.4 Attention: The Assessment Mirror

7.4.1 Educational Attention (Learning Outcomes)

Contemporary education forces student attention through predetermined “learning outcomes”: specific, measurable, achievable, relevant, time-bound (SMART) objectives that determine what matters. Everything else becomes noise to be filtered.

A literature course that once explored infinite interpretations of Hamlet now optimizes for:

- “Identify three themes in Acts 1-3” (measurable)
- “Compare two critical interpretations” (assessable)
- “Write 2,000-word analysis” (quantifiable)

The student’s attention is forcibly directed to what will be tested. Wonder, curiosity, and tangential insight (the seeds of genuine understanding) are filtered as inefficiencies. The system implements attention mechanisms that eliminate everything except what optimizes assessment scores.

7.4.2 Computational Attention

The transformer's attention mechanism implements the mathematical formula:

$$\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right) V \quad (3)$$

Where:

- Q (Query): What we're looking for
- K (Key): What's available to match
- V (Value): What gets retrieved
- Softmax: Forces focus on highest scores

This IS the standardized examination:

- Query: Test question
- Key: Possible answers in memory
- Value: Creditable responses
- Softmax: Grade curve forcing discrimination

The mechanism forces focus on predetermined patterns while systematically filtering everything else. Just as students learn to attend only to what affects grades, transformers attend only to what affects loss functions.

7.5 The Training Parallel

The LLM training process mirrors the human educational timeline with disturbing precision:

7.5.1 Phase 1: Pre-training (Comprehensive Absorption)

- **LLM:** Consumes massive text corpus without judgment or discrimination
- **Medieval Education:** Seven years of trivium/quadrivium, absorbing all knowledge domains
- **Energy:** Maximum investment, no immediate output expected
- **Result:** Broad capability foundation

7.5.2 Phase 2: Fine-tuning (Specialization)

- **LLM:** Narrow training on specific domains/tasks
- **Bologna Bachelor:** Three years focused specialization
- **Energy:** Reduced investment, targeted output
- **Result:** Domain-specific performance

7.5.3 Phase 3: RLHF (Compliance Training)

- **LLM:** Reinforcement learning from human feedback to eliminate “undesirable” outputs
- **Contemporary Assessment:** Continuous testing to ensure compliance with expected responses
- **Energy:** Minimal investment, maximum control
- **Result:** Predictable, “safe” outputs

Both progressions move from energy-intensive comprehensiveness toward energy-minimal compliance. Both sacrifice capability for control.

7.6 The Recursive Feast

The most horrifying revelation: LLMs now train on text produced by humans who were trained to think like machines. The recursive loop accelerates:

1. **Generation 1:** Humans trained to process information algorithmically
2. **Generation 2:** Machines trained on algorithmically-processed human outputs
3. **Generation 3:** Humans learning from machines trained on mechanized human thought
4. **Generation 4:** Machines learning from humans who learned from machines...

Each iteration loses additional depth. The LLM trained on academic papers written by scholars who were trained to write for impact factors produces text optimized for...impact factors. The system converges toward perfect emptiness: maximum optimization, minimum meaning.

OpenAI’s GPT models demonstrate this progression:

- GPT-2 (2019): Trained on wild internet text, occasional brilliance amid chaos

- GPT-3 (2020): Trained on curated text, more consistent but less surprising
- GPT-4 (2023): Trained on refined data plus human feedback, reliable but predictable
- Future models: Training on AI-generated text, approaching semantic heat death

7.7 The Thermodynamic Proof

The energy requirements reveal the fundamental difference:

7.7.1 Human Cognition (Historical)

- Formation: 20 years continuous biological energy investment
- Maintenance: Lifetime energy requirement for neuroplasticity
- Adaptation: Constant energy for contextual learning
- Creativity: High-energy states enabling novel connections

7.7.2 LLM “Intelligence”

- Training: One-time massive energy expenditure (1,287 MWh for GPT-3)
- Inference: Minimal energy for pattern matching
- Adaptation: Zero (frozen weights after training)
- Creativity: None (statistical recombination only)

Humans were negative entropy systems: local reversals of thermodynamic law through continuous energy investment. LLMs are entropy crystals: frozen patterns that can only decay. We’ve trained humans to be more like LLMs: front-loaded training creating static patterns rather than continuous adaptive growth.

7.8 The Perfect Mirror

The AI mirror reveals our self-portrait: we see in LLMs exactly what we’ve become. They process tokens because we process credits. They embed in vector spaces because we embed in competency frameworks. They attend selectively because we assess selectively. They optimize for loss functions because we optimize for grades.

The machines aren’t becoming conscious; we’re becoming mechanical. The convergence point isn’t artificial general intelligence but biological specific processing. We meet

our creations halfway, in the diminished middle where neither human wisdom nor machine efficiency exists, only the automated processing of pre-processed patterns.

The mirror is perfect because we ground both sides to match. LLMs are successfully replacing human cognitive work not because they’ve achieved human capability but because we’ve reduced human capability to what machines can replicate.

We trained ourselves for replacement. The machines simply arrived to occupy the positions we’d prepared.

8 Reconstruction Principles: Building Cognitive Sovereignty Within Thermodynamic Constraints

8.1 The Epistemological Prerequisites

Before proposing reconstruction pathways, we must acknowledge a fundamental paradox: those seeking to rebuild cognitive sovereignty are themselves products of the vectorization they seek to escape. The map is never the territory ([Korzybski, 1933](#)), and our maps were drawn with vectorized tools.

Contemporary decision science reveals the illusion of individual rationality. As [Kahneman \(2011\)](#) demonstrated, human decision-making operates through systematic biases rather than rational calculation. More fundamentally, distributed cognition theory ([Hutchins, 1995](#)) establishes that thinking occurs not within individual minds but across socio-technical systems. What we experience as “our” decisions emerge from complex interactions between cultural values, institutional practices, environmental constraints, and collective sense-making processes.

Those educated in STEM disciplines (precisely those positioned to understand technical systems) are most deeply conditioned by bounded rationality ([Simon, 1991](#)). The analytical tools revealing the vectorization problem are themselves products of vectorized thinking. We are attempting to examine the lens through which we see using that same lens: [Hofstadter \(1979\)](#) strange loop made manifest.

Critical retrospection of any recent non-trivial decision reveals this embedding. The choice to pursue alternative educational approaches, implement new organizational structures, or resist AI integration emerges not from individual cognition but from:

- Cultural layer: Prevailing beliefs about knowledge and value ([Schein, 1985](#))
- Institutional layer: Organizational structures and incentive systems ([DiMaggio and Powell, 1983](#))
- Social layer: Peer networks and professional communities ([Granovetter, 1985](#))
- Individual layer: Personal history and embodied experience ([Bourdieu, 1990](#))

The individual functions as one node in a distributed processing network: “the literal neuron of a bigger brain.” This recognition demands epistemic humility. We cannot stand outside the system to reconstruct it. As [Maturana and Varela \(1987\)](#) establish, we are “structurally coupled” to the environment we seek to transform.

8.2 The Cynefin Diagnostic Framework

Snowden’s Cynefin framework ([Snowden and Boone, 2007](#)) provides the essential diagnostic tool for understanding where reconstruction is both necessary and possible. The framework distinguishes five domains, each requiring different cognitive architectures:

Clear Domain (Known knowns)

- Best practices apply
- Sense $\rightarrow$ Categorize $\rightarrow$ Respond
- Vectors excel here: procedural, repeatable, optimizable
- High AI digestibility

Complicated Domain (Known unknowns)

- Good practices exist through expertise
- Sense $\rightarrow$ Analyze $\rightarrow$ Respond
- Vectors vulnerable: expertise reducible to procedures
- Moderate AI digestibility

Complex Domain (Unknown unknowns)

- Emergent practices required
- Probe $\rightarrow$ Sense $\rightarrow$ Respond
- Spheres essential: no predetermined responses possible
- Low AI digestibility

Chaotic Domain (No cause-effect relationship)

- Novel practices needed
- Act $\rightarrow$ Sense $\rightarrow$ Respond
- Spheres critical: immediate embodied response required

- Near-zero AI digestibility

Confused Domain (Unclear which domain applies)

- Most dangerous state
- Requires meta-cognitive capacity to diagnose
- Sphere capability: domain recognition itself
- Cannot be automated

The critical insight: organizations systematically misclassify complex problems as complicated, applying “best practices” where emergent approaches are needed. [Ford et al. \(2024\)](#) quantified this: 341 “Simple” problems were actually Complicated, 437 “Complicated” were Complex, creating “catastrophic failure” when complicated approaches fail in complex domains.

8.3 Thermodynamic Constraints on Reconstruction

Reconstruction faces inescapable thermodynamic constraints. The Second Law doesn’t negotiate. Energy invested in cognitive development cannot be recovered from entropic systems, and new investment requires energy sources increasingly consumed by existing structures.

The Energy Equation:

$$\text{Cognitive Sovereignty} = \frac{\text{Energy Invested}}{\text{Time}} \times \text{Resistance to Extraction} \quad (4)$$

Where: Energy Investment > Entropy Rate

Historical benchmarks reveal the required investment scale:

- Sophia (theoretical wisdom): 20+ years sustained investment
- Phronesis (practical wisdom): Lifetime daily practice
- Techne (craft knowledge): 10,000+ hours minimum ([Ericsson et al., 1993](#))
- Metis (adaptive cunning): Constant challenge exposure
- Nous (intuitive insight): Unknown threshold, possibly unreachable

Modern constraints make these investments nearly impossible:

- Work demands: 40-60 hours/week vectorized activity
- Economic pressure: Continuous productivity requirements

- Attention economy: Constant extraction attempts
- Social expectations: Optimization for measurable outcomes

Realistic reconstruction must acknowledge these constraints while creating protected spaces for energy investment.

8.4 Individual Reconstruction Protocols

8.4.1 Phase 1: Diagnostic Assessment (Month 1)

Current State Analysis

- Map daily activities to Cynefin domains
- Identify vector lock-in patterns
- Calculate actual discretionary time/energy
- Assess extraction vulnerability

Sphere Potential Mapping

- Identify existing cross-domain connections
- Recognize latent capacities from pre-vectorization
- Map curiosity patterns that resist optimization
- Locate energy reserves for investment

Realistic Timeline Acceptance

Accept thermodynamic reality:

- No shortcuts exist (physics doesn't negotiate)
- Decades required for sphere development
- Starting now means barely enough time
- Partial development better than none

8.4.2 Phase 2: Foundation Building (Months 2-6)

Cross-Domain Exposure

- Read deliberately outside primary field (minimum 2 hours/week)
- Attend events in unfamiliar domains
- Engage communities with different knowledge traditions
- Follow curiosity without optimization goals

Pattern Recognition Development

- Maintain synthesis journal for cross-domain insights
- Practice analogical thinking between disparate fields
- Build personal pattern library
- Resist premature categorization

Embodied Practice Initiation

Essential for developing non-extractable knowledge:

- Physical craft (woodworking, pottery, gardening)
- Movement practice (martial arts, dance, climbing)
- Musical instrument learning
- Somatic awareness development

8.4.3 Phase 3: Cultivation (Months 6-24)

Deliberate Integration

- Create projects requiring multiple domain knowledge
- Write integrative analyses crossing boundaries
- Teach others using cross-domain metaphors
- Build synthesis as default mode

Complexity Navigation Practice

- Seek problems without predetermined solutions
- Practice probe-sense-respond in safe contexts

- Embrace emergence and uncertainty
- Document pattern recognition development

Community Building

Sphere development requires collective support:

- Find others attempting reconstruction
- Create learning partnerships
- Share failures and insights
- Build counter-cultural spaces

8.5 Organizational Reconstruction Architecture

8.5.1 From Silos to Sphere Teams

Replace functional specialization with cross-domain capacity:

- Pilot teams mixing 3-5 different expertise domains
- Complex challenge focus rather than efficiency optimization
- Protected learning time (minimum 20% non-productive)
- Success measured by adaptation not standardization

8.5.2 From Best Practices to Emergent Practices

Shift from standardization to contextual response:

- Train all staff in Cynefin framework
- Create “probe” budgets for complex challenges
- Document emergent solutions without standardizing
- Celebrate contextual wisdom over universal procedures

8.5.3 From Extraction to Investment

Reverse the energy flow:

- Implement genuine apprenticeship programs (3+ years)
- Create internal “universities” for cross-domain learning
- Protect thinking time from productivity metrics
- Measure knowledge development not just application

8.6 Pragmatic Implementation Protocols

8.6.1 Week 1 Actions

- Complete Cynefin self-assessment for current work
- Identify one complex challenge being treated as complicated
- Block 4 hours for cross-domain exploration
- Start synthesis journal

8.6.2 Month 1 Targets

- Read two books from unrelated fields
- Attend one event outside professional domain
- Begin one embodied practice
- Connect with three people from different backgrounds

8.6.3 Quarter 1 Objectives

- Establish 10 hours/week protected development time
- Complete initial probe-sense-respond project
- Build learning partnership or join community
- Document energy investment patterns

8.6.4 Year 1 Goals

- Develop beginning competence in one non-professional domain
- Navigate five complex challenges using emergent practices
- Build sustainable energy investment habits
- Create or contribute to sphere development community

8.7 Critical Success Factors and Failure Patterns

Success Requirements:

- Accept decades-long timeline (thermodynamic reality)
- Protect energy investment despite productivity pressure
- Build community support (individual reconstruction impossible)
- Measure progress in capability not credentials
- Maintain sovereignty despite extraction attempts

Common Failure Patterns:

- Treating as temporary project rather than permanent practice
- Attempting alone without community support
- Measuring by vectorized metrics (efficiency, speed)
- Expecting linear progress in complex domain
- Underestimating energy investment required

The Fundamental Choice:

Reconstruction isn't optimization or self-improvement. It's choosing cognitive sovereignty over efficiency, wisdom over productivity, and decades of investment over immediate returns. The choice is binary: invest the energy or accept dissolution. Physics doesn't negotiate.

9 Implications: When Physics Meets Mythology

9.1 The Epistemological Collapse: How We Built Truth from Lies

The most devastating revelation isn't that transformation initiatives fail; it's how we created the knowledge about their failure. [Hughes \(2011\)](#) traces the genealogy of the "70% failure rate" and discovers... nothing. [Beer and Nohria \(2000\)](#) stated it as "brutal fact" without evidence. [Kotter \(2008\)](#) cited it as accepted truth without references. Through what Hughes calls "unconscious collusion," fiction became fact through citation networks.

Then reality arrived. [Bain & Company \(2024\)](#): 88% failure rate with actual data from 400+ executives. [Boston Consulting Group \(2024\)](#): 74% failure rate from analyzing

1,700+ transformations. [Ford et al. \(2024\)](#): 1,205 non-conformance reports showing systematic complexity misclassification. The mythology was accidentally correct, but for entirely wrong reasons.

This is our entire thesis in microcosm: We create simplified models (70% failure myth), they become institutionalized truth (cited thousands of times), then reality proves even worse than the fiction (88% actual failure), but we can't see why because we're trapped in our own simplification.

Kotter's confession is breathtaking: "I neither drew examples nor major ideas from any published source, except my own writing" ([Hughes, 2015](#)). The most influential change management framework of the past 30 years (8 steps taught in every MBA program, implemented in thousands of organizations) has ZERO empirical foundation. Built pre-internet, surviving post-Google, thriving in the age of AI. Pure mythology dressed as methodology.

[McLaren et al. \(2023\)](#) identify the fatal contradiction: Kotter admits people prefer status quo over uncertainty, then demands leaders make "the current situation look more dangerous than launching into the unknown." Working against human psychology while claiming to manage human change. The model defeats itself.

9.2 The German Proof: When Industry Knows What Academia Denies

TU Dresden engineering faculty state it plainly: "The university Bachelor's degree in six semesters does not lead to a professionally qualifying degree... essential components must be omitted... This damages the foundation of engineering education" ([Odenbach and Krauthäuser, 2015](#)).

German industry votes with wallets. Collective bargaining agreements assign different wage groups to Diplom vs Bachelor/Master holders. Master's graduates earn wage premiums "even relative to those with more work experience" ([Wieschke et al., 2020](#)). The TU9 consortium (educating 47% of German engineers) formally requested the right to award "Diplom-Ingenieur" alongside Master's degrees because employers know the difference.

[Kaiser and Schröder \(2022\)](#) quantify what's missing: "Systems Thinking, collaboration and communication are not explicitly addressed" in modularized engineering education. The very capabilities that distinguish engineers from algorithms (seeing wholes, navigating emergence, integrating across domains) systematically eliminated by Bologna optimization.

VDI President Ungeheuer's lament: "We sacrificed the Diplom-Ingenieur with heavy hearts for a greater goal, namely international connectivity." They traded thermodynamic reality for bureaucratic compatibility. Physics doesn't recognize international agreements.

9.3 The Micro-Credential Delusion: Approaching Absolute Zero

[Lumina Foundation \(2025\)](#) celebrates: 96% of employers say micro-credentials strengthen applications. 90% offer 10-15% higher starting salaries. 87% hired credential holders last year. The numbers look magnificent.

Then examine actual performance. [Ha et al. \(2022\)](#) find only one negative outcome study among 14 effectiveness assessments: because nobody measures long-term degradation. [Gauthier \(2020\)](#) reveals the truth: employers prefer micro-credentials not because they indicate capability but because degrees have already become meaningless. When traditional credentials fail to communicate competence, granular badges seem like improvement. Racing toward zero, celebrating each milestone of decline.

[Joshi \(2019\)](#) provides the thermodynamic proof: bootcamps help non-technical graduates enter tech but provide “minimal benefit for those already holding technical degrees”; you can’t add knowledge to an empty container, but you can’t accelerate what doesn’t exist. The energy investment was already absent.

9.4 The Complexity Catastrophe: Treating Cancer with Band-Aids

[Kempermann \(2017\)](#) on biomedicine: “Complex problems are often treated as if they were not more than the complicated sum of solvable sub-problems...this is not correct [and has] dangerous consequences, especially in clinical contexts.” People die from category errors.

[Ford et al. \(2024\)](#) quantify organizational blindness: analyzing 1,205 quality problems in a £1.45 billion megaproject:

- 341 “Simple” problems were actually Complicated
- 437 “Complicated” problems were actually Complex
- Systematic misclassification led to repeated failures

This isn’t incompetence; it’s thermodynamic inevitability. Organizations optimize for efficiency (simple/complicated domains) while reality presents complexity. As [Alexander et al. \(2018\)](#) demonstrate, performance measurement systems assume predictability that doesn’t exist. We measure what we can control, ignore what we can’t, then act surprised when unmeasured reality destroys measured fantasy.

9.5 The AI Mirror Confirms: We Built Our Replacements

IBM’s saga crystallizes everything: Replace 8,000 HR workers with AI. Discover AI handles 94% of routine tasks brilliantly. Then discover the 6% requiring “empathy, nuance,

trust” makes the system non-functional. Forced to rehire because [Collins \(2010\)](#) Collective Tacit Knowledge can’t be extracted.

Duolingo’s cliff dive: Fire 100+ curriculum experts who understood language as cultural embodiment. Replace with GPT trained on their extracted patterns. Users with 1,131-day streaks quit in protest. Content becomes “repetitive, robotic” lacking “playful tone” and “cultural nuance.” The sphere-holders created what vectors now fail to maintain.

Microsoft’s Althoff: “\$500 million saved using AI in 2024” from 15,000 layoffs. But as critics note: “we don’t know what metrics the company uses...[it could be] just the combined salaries of the thousands of people the company laid off.” Quantifying extraction, ignoring depletion. Thermodynamic debt accumulating, payment deferred.

9.6 The Binary Future: Physics Doesn’t Negotiate

The implications converge on a single, non-negotiable reality: Systems violating thermodynamic law will collapse. Not might, not could: will. The Second Law doesn’t read quarterly reports or respect international agreements.

For Education: Every micro-credential is acceleration toward heat death. The Lumina Foundation celebrating 96% employer approval while cognitive capacity degrades is like celebrating how fast we’re approaching the cliff. TU Dresden maintaining five-year integrated Diplom programs while others fragment into badges shows that resistance remains possible, but only through energy investment.

For Organizations: The 88% transformation failure rate ([Bain & Company, 2024](#)) isn’t about poor execution; it’s about thermodynamic violation. You cannot reorganize entropy away. Real transformation requires sphere development: years of investment, protected learning time, accepting efficiency loss for capability gain. McKinsey selling three-month transformations is literally selling perpetual motion machines.

For Individuals: The choice is binary. Either invest energy in sphere development (accepting decade-long timelines, resisting optimization pressure, building cross-domain capacity) or accept vectorization and eventual replacement. There is no middle path because physics doesn’t compromise.

9.7 The Kotter Memorial: A Case Study in Civilizational Failure

John Kotter (Harvard Business School professor emeritus, author of 20 books, consultant to Fortune 500 companies) built the most influential change management framework of the past 30 years on literally nothing. No empirical evidence. No theoretical foundation. No external sources. Just personal experience marketed as universal truth.

That this is possible (that academia's quality control failed this completely, that thousands of organizations implemented fantasy as methodology, that MBA programs still teach it despite Hughes' devastating critique) proves our thesis absolutely. We don't verify, we don't validate, we don't even check sources. We accept simplified models that feel true, cite them into existence, then act shocked when reality refuses to comply.

Kotter succeeded because he offered what vectors want: eight simple steps, linear progression, measurable stages, the illusion of control. He failed because complexity doesn't care about our simplifications. The 88% failure rate isn't despite following Kotter; it's because we followed Kotter.

9.8 The Thermodynamic Reckoning

Every institution operating on extraction without investment faces the same endpoint. Every optimization that reduces energy input accelerates entropy. Every simplification that ignores complexity ensures catastrophic failure. Every vectorization that destroys spheres guarantees replacement by machines that process vectors more efficiently.

The German engineers know this, maintaining their Diplom against EU pressure. A few holdout institutions preserve sphere development. Individual practitioners build local negentropy bubbles. But the overall trajectory is clear: thermodynamic collapse accelerating.

We trained humans to think like machines, documented the training exhaustively, then built machines that think like we trained humans to think. Now we discover that humans trained to think like machines are inferior to actual machines, while humans who think like humans are increasingly rare, increasingly valuable, and increasingly impossible to develop within our current systems.

The feast hasn't begun; it's concluding. The appetizers (routine work) have been consumed. The main course (professional work) is being served. Only the indigestible spheres remain.

Build them or be consumed. Physics doesn't negotiate.

10 Conclusion: The Thermodynamic Reckoning

The evidence converges on an inescapable conclusion: we systematically transformed human cognition into something replaceable, documented the process exhaustively, then built the replacements. This wasn't conspiracy but optimization: each local decision rational, the collective outcome thermodynamically inevitable.

I've given you the physics. Now I'll give you my life as proof that the physics can be navigated.

I escaped academic vectorization—walked away from a PhD for the internet in 1995.

I escaped technical vectorization—left CTO roles for consulting. I escaped corporate vectorization—bought a Unimog and drove to South America when the optimization trap was killing me. I escaped geographic vectorization—returned from American turbo-capitalism to German complexity. Six vector traps, six escapes. Each required energy investment. Each required accepting inefficiency. Each required choosing sovereignty over optimization.

These aren't romantic gestures. They're thermodynamic proof: the Second Law describes tendency, not destiny. Local negentropy bubbles can be built. Spheres can be maintained. But only through continuous energy investment that our systems have been designed to prevent.

The historical arc is clear. From ancient academies investing decades in multidimensional wisdom to micro-credentials measured in hours, we've followed an exponential decay curve toward cognitive zero. The Bologna Process didn't just modularize education; it modularized minds. Organizations didn't just optimize processes; they optimized away judgment itself. We didn't just build AI; we first rebuilt humans to think like the AI we would build.

The thermodynamic framework reveals why reconstruction is difficult but not impossible. Knowledge isn't information—it's a high-energy state requiring continuous investment to maintain. Every efficiency gain that reduces energy input accelerates entropy. Every standardization that enables measurement destroys the unmeasurable. Every vector that replaces a sphere makes the system more optimized and more fragile, more efficient and less adaptable.

The contemporary failures—42% AI abandonment rate, 88% transformation failure rate, forced rehiring after AI replacement—aren't implementation problems. They're physics problems. You cannot extract what was never invested. You cannot automate what you don't understand. You cannot replace spheres with vectors and expect the system to navigate complexity.

The confession literature provides the most damning evidence, and the clearest hope. Educational psychologists documented their standardization techniques—which means the techniques can be identified and reversed. Management consultants celebrated their extraction methods—which means the methods can be recognized and resisted. Platform designers published their manipulation strategies—which means the strategies can be seen and circumvented. They confess openly because they see no crime. Their confession is our map.

The German engineers holding onto their Diplom, TU Dresden maintaining integrated programs, individual practitioners building local negentropy bubbles: these aren't romantic holdouts but thermodynamic realists. They understand what the efficiency optimizers refuse to see.

For individuals: entropy isn't waiting for your decision. While you contemplate op-

tions, the Second Law is already working. The question isn't whether to invest in cognitive development but whether you've already started. Every hour of cross-domain synthesis, every practice session building embodied skill, every deliberate resistance to optimization pressure—these are the energy inputs that maintain sovereignty. There is no pause button.

For organizations: the 88% transformation failure rate will approach 100% as vectorized knowledge proves insufficient for complexity navigation. Only organizations investing in genuine capability development—years not quarters, learning not training, emergence not planning—have any chance of navigating what's coming. Physics doesn't read business plans.

For civilization: we face a closing window. The last generation that experienced comprehensive education is retiring. The last institutions maintaining integrated development are fragmenting into modules. Once the knowledge of how to develop spheres is lost, reconstruction becomes archaeology rather than pedagogy.

The feast hasn't begun; it's concluding. We prepared the meal ourselves, seasoned it with our own standardization, served it on the plate of our own optimization. But some of us refused to eat. Some of us built local kitchens. Some of us remembered what food tasted like before it was processed.

That memory is the seed. The thermodynamics are simple: invest energy or watch entropy work. The window is closing, but it hasn't closed.

Physics doesn't negotiate. Neither should we.

References

- Aaker, D. A. (2008). Spanning silos: The new cmo imperative. *Harvard Business Review*, 86(5):68–76.
- Alexander, A. et al. (2018). A decision theory perspective on complexity in performance measurement and management. *International Journal of Operations & Production Management*, 38(11):2214–2244.
- Althoff, J. (2024). Internal presentation on ai productivity gains.
- Alvesson, M. (2013). *The Triumph of Emptiness: Consumption, Higher Education, and Work Organization*. Oxford University Press.
- Ashkenas, R. (2015). Why we still can't break down the silos. *Harvard Business Review Digital Articles*, pages 2–5.
- Bain & Company (2024). 88% of business transformations fail to achieve their original ambitions. Technical report, Bain & Company.

- Beer, M. and Nohria, N. (2000). Cracking the code of change. *Harvard Business Review*, 78(3):133–141.
- Beller, M. (1996). The conceptual and the anecdotal history of science. *Studies in History and Philosophy of Science*, 27:167–194.
- Boston Consulting Group (2024). Five truths and one lie about corporate transformation. Technical report, BCG.
- Bourdieu, P. (1990). *The Logic of Practice*. Stanford University Press.
- Bratianu, C. (2020). Knowledge dynamics: A thermodynamics approach. *Kybernetes*, 49(1):6–21.
- Bratianu, C. and Bejinaru, R. (2020). Knowledge dynamics: A thermodynamics approach. *Kybernetes*, 49(2):22–30.
- Brillouin, L. (1956). *Science and Information Theory*. Academic Press, New York.
- Brynjolfsson, E. and McAfee, A. (2014). *The Second Machine Age: Work, Progress, and Prosperity in a Time of Brilliant Technologies*. W. W. Norton & Company, New York.
- Burke, P. (2020). *The Polymath: A Cultural History from Leonardo da Vinci to Susan Sontag*. Yale University Press, New Haven, CT.
- Collins, H. (2010). *Tacit and Explicit Knowledge*. University of Chicago Press, Chicago, IL.
- Collins, J. (2022). The 1/n portfolio versus the optimal strategy: Does a simple heuristic beat the nobel prize winner’s solution? Retrieved January 20, 2022.
- De la Croix, D., Doepke, M., and Mokyr, J. (2018). Clans, guilds, and markets: Apprenticeship institutions and growth in the preindustrial economy. *The Quarterly Journal of Economics*, 133(1):1–70.
- DeLanda, M. (2006). *A New Philosophy of Society: Assemblage Theory and Social Complexity*. Continuum.
- Deleuze, G. and Guattari, F. (1987). *A Thousand Plateaus: Capitalism and Schizophrenia*. University of Minnesota Press.
- DeMiguel, V., Garlappi, L., and Uppal, R. (2009). Optimal versus naive diversification: How inefficient is the 1/N portfolio strategy? *Review of Financial Studies*, 22(5):1915–1953.

- DiMaggio, P. J. and Powell, W. W. (1983). The iron cage revisited: Institutional isomorphism and collective rationality in organizational fields. *American Sociological Review*, 48:147–160.
- Dreyfus, H. L. (1979). *What Computers Can't Do: The Limits of Artificial Intelligence*. MIT Press.
- Edmondson, A. C. (2012). *Teaming: How Organizations Learn, Innovate, and Compete in the Knowledge Economy*. Jossey-Bass.
- Endsley, M. R. (1995). Toward a theory of situation awareness in dynamic systems. *Human Factors*, 37(1):32–64.
- Ericsson, K. A. (2006). The influence of experience and deliberate practice on the development of superior expert performance. *Cambridge Handbook of Expertise and Expert Performance*, pages 683–703.
- Ericsson, K. A. et al. (1993). The role of deliberate practice in the acquisition of expert performance. *Psychological Review*, 100(3):363–406.
- Erkizan, H. N. (1997). *Energeia, nous and non-discursive thinking in Aristotle*. PhD thesis, University of Bristol.
- Faber, W. (2022). Harry markowitz: An equal-weight investor? Alpha Architect Research, May 2.
- Feynman, R. P. (1974). Cargo cult science. *Engineering and Science*, 37(7):10–13.
- Ford, G. et al. (2024). Simplifying complexity? on quality decision-making and non-conformance outcomes of megaprojects. *IEEE Transactions on Engineering Management*, 71:5443–5454.
- Gadamer, H.-G. (1960). *Truth and Method*. Continuum. English translation 1975.
- Galbraith, J. R. (1973). *Designing Complex Organizations*. Addison-Wesley, Reading, MA.
- Gauthier, T. (2020). The value of microcredentials: The employer's perspective. *The Journal of Competency-Based Education*, 5(2):e01209.
- Georgescu-Roegen, N. (1971). *The Entropy Law and the Economic Process*. Harvard University Press.
- Gigerenzer, G. (2014). *Risk Savvy: How to Make Good Decisions*. Viking, New York, NY.

- Gigerenzer, G. and Selten, R. (2001). *Bounded Rationality: The Adaptive Toolbox*. MIT Press.
- Gleeson, J. (2021). The european credit transfer system (ects) from bologna to the market. *Policy Futures in Education*, 19(4):492–508.
- Granovetter, M. (1985). Economic action and social structure: The problem of embeddedness. *American Journal of Sociology*, 91:481–510.
- Ha, N. T. N. et al. (2022). A systematic literature review of micro-credentials in higher education: A non-zero-sum game. *Higher Education Research and Development*, 42(4):888–906.
- Hackman, J. R. (2002). Leading teams: Setting the stage for great performances. *Harvard Business School Press*.
- History Hit (2025). What did european universities teach during the middle ages?
- Hoel, E. (2025). Polymaths are back from the dead. *The Intrinsic Perspective*.
- Hofstadter, D. R. (1979). *Gödel, Escher, Bach: An Eternal Golden Braid*. Basic Books.
- Hughes, M. (2011). Do 70 per cent of all organizational change initiatives really fail? *Journal of Change Management*, 11(4):451–464.
- Hughes, M. (2015). Leading changes: Why transformation explanations fail. *Leadership*, 11(4):435–460.
- Hutchins, E. (1995). *Cognition in the Wild*. MIT Press.
- Jamadar, S. D. (2025). The metabolic costs of cognition. *Trends in Cognitive Sciences*, 29(1):45–59.
- Joshi, S. (2019). *Technical job placement success of coding bootcamp graduates*. PhD thesis, University of Pennsylvania.
- Kahneman, D. (2011). *Thinking, Fast and Slow*. Farrar, Straus and Giroux.
- Kahneman, D. and Klein, G. (2009). Conditions for intuitive expertise: A failure to disagree. *American Psychologist*, 64(6):515–526.
- Kaiser, L. and Schröder, E. (2022). Systems engineering in university education with focus on systems thinking. *Proceedings of Design Society*, 3:2871–2880.
- Kempermann, G. (2017). Cynefin as reference framework to facilitate insight and decision-making in complex contexts of biomedical research. *Frontiers in Neuroscience*, 11:634.

- Klein, G. A. (1993). *A recognition-primed decision (RPD) model of rapid decision making*. Ablex Publishing.
- Korzybski, A. (1933). *Science and Sanity: An Introduction to Non-Aristotelian Systems and General Semantics*. Institute of General Semantics.
- Kotter, J. P. (2008). *A Sense of Urgency*. Harvard Business Press.
- Krishna, A. (2024). Ibm ceo confirms ai replacement of hundreds of hr staff.
- Kroher, M. et al. (2021). Social inequalities in higher education after bologna. *European Journal of Higher Education*, 11:261–278.
- Kuhn, T. S. (1962). *The Structure of Scientific Revolutions*. University of Chicago Press.
- Landauer, R. (1961). Irreversibility and heat generation in the computing process. *IBM Journal of Research and Development*, 5(3):183–191.
- Lowenstein, R. (2000). *When Genius Failed: The Rise and Fall of Long-Term Capital Management*. Random House.
- Lumina Foundation (2025). Micro-credentials impact report 2025: Insights from students and employers. Technical report, Lumina Foundation.
- Lweis, F. I. et al. (2022). Medicine, value, and knowledge in the veterinary clinic. *Sociology of Health & Illness*. PMC8964002.
- Maturana, H. R. and Varela, F. J. (1987). *The Tree of Knowledge: The Biological Roots of Human Understanding*. Shambhala.
- McKinsey & Company (2025). Seizing the agentic ai advantage. Technical report, McKinsey & Company.
- McLaren, T. A. S. et al. (2023). Why vilifying the status quo can derail a change effort: Kotter’s contradiction, and theory adaptation. *Journal of Change Management*, 23(1):93–111.
- McNamee, D. (2019). Cargo cult analytics: The danger of unexamined metrics. *Harvard Business Review*. Online article.
- Medieval Buildings (n.d.). Carpenters’ knowledge.
- Morgan, M. G. (2014). Use (and abuse) of expert elicitation in support of decision making for public policy. *Proceedings of the National Academy of Sciences*, 111(20):7176–7184.
- Moulier-Boutang, Y. (2007). *Le capitalisme cognitif: La nouvelle grande transformation*. Editions Amsterdam.

- Mouzala, M. G., editor (2024). *Cognition in Ancient Greek Philosophy and its Reception: Interdisciplinary Approaches*. Academia Verlag.
- Mulder, M., Gulikers, J., Biemans, H., and Wesselink, R. (2007). The new competence concept in higher education: Error or enrichment? *Journal of European Industrial Training*, 31(8):755–770.
- Nonaka, I. and Konno, N. (1998). The concept of “ba”: Building a foundation for knowledge creation. *California Management Review*, 40(3):40–54.
- Nonaka, I. and Takeuchi, H. (1995). *The Knowledge-Creating Company: How Japanese Companies Create the Dynamics of Innovation*. Oxford University Press, New York.
- Odenbach, S. and Krauthäuser, H. G. (2015). Mehr als ein akademischer grad: Plädoyer für das diplom in den ingenieurwissenschaften. *Forschung & Lehre*, 6:450–451.
- Peña, A. (2010). The dreyfus model of clinical problem-solving skills: A critical perspective. *Medical Education Online*, 15(1):4846.
- Polanyi, M. (1966). *The Tacit Dimension*. Doubleday, Garden City, NY.
- Powell, W. W. and Snellman, K. (2004). The knowledge economy. *Annual Review of Sociology*, 30:199–220.
- Raichle, M. E. and Gusnard, D. A. (2002). Appraising the brain’s energy budget. *Proceedings of the National Academy of Sciences*, 99(16):10237–10239.
- Rinta-Kahila, T. et al. (2023). Algorithmic decision-making and system destructiveness: A case of automatic debt recovery. *European Journal of Information Systems*, 32(3):45–62.
- Robinson, A. (2006). *The Last Man Who Knew Everything: Thomas Young, the Anonymous Polymath Who Proved Newton Wrong*. Pi Press.
- Ross University School of Veterinary Medicine (n.d.). Is vet school harder than medical school?
- Schein, E. H. (1985). *Organizational Culture and Leadership*. Jossey-Bass.
- Schrödinger, E. (1944). *What is Life? The Physical Aspect of the Living Cell*. Cambridge University Press.
- Senge, P. M. (1990). *The Fifth Discipline: The Art and Practice of the Learning Organization*. Doubleday.

- Shannon, C. E. (1948). A mathematical theory of communication. *Bell System Technical Journal*, 27(3):379–423.
- Simon, H. A. (1991). Bounded rationality and organizational learning. *Organization Science*, 2(1):125–134.
- Smil, V. (2017). *Energy and Civilization: A History*. MIT Press, Cambridge, MA.
- Snowden, D. J. and Boone, M. E. (2007). A leader’s framework for decision making. *Harvard Business Review*, 85(11):68–76.
- S&P Global Market Intelligence (2025). Ai experiences rapid adoption, but with mixed outcomes. Technical report, S&P Global.
- Srnicek, N. (2017). *Platform Capitalism*. Polity.
- Stonier, T. (1996). Information as a basic property of the universe. *Biosystems*, 38(2-3):135–140.
- Susskind, R. and Susskind, D. (2020). *The Future of the Professions: How Technology Will Transform the Work of Human Experts*. Oxford University Press, Oxford, updated edition edition.
- Taleb, N. N. (2001). *Fooled by Randomness: The Hidden Role of Chance in Life and in the Markets*. Random House.
- Taleb, N. N. (2007). *The Black Swan: The Impact of the Highly Improbable*. Random House.
- Tetlock, P. E. (2005). *Expert Political Judgment: How Good Is It? How Can We Know?* Princeton University Press.
- Vandeweerd, J.-M., Kirschvink, N., Clegg, P., Vandenput, S., Gustin, P., and Saegerman, C. (2013). Clinical reasoning and case-based decision making: The fundamental challenge to veterinary educators. *Journal of Veterinary Medical Education*, 40(3):200–209.
- Vaswani, A. et al. (2017). Attention is all you need. *Advances in Neural Information Processing Systems*, 30:5998–6008.
- Vercellone, C. (2007). From formal subsumption to general intellect. *Historical Materialism*, 15:13–36.
- von Ahn, L. (2024). Making duolingo ai-first. Internal memo published on LinkedIn.
- Watkins, M. D. (2013). *The First 90 Days: Proven Strategies for Getting Up to Speed Faster and Smarter*. Harvard Business Review Press, updated and expanded edition.

- Wiehler, A. et al. (2022). A neuro-metabolic account of why daylong cognitive work alters economic preferences. *Current Biology*, 32(16):3564–3575.
- Wieschke, J., Reimer, M., and Falk, S. (2020). Labour market returns of bachelor’s and master’s degrees in germany. *Beiträge zur Hochschulforschung*, 42(3):56–78.
- Wikipedia (2002). Medieval university. Accessed 2025.
- Zuboff, S. (2019). *The Age of Surveillance Capitalism*. PublicAffairs.